

Process Over Magic: Beyond Vibe Coding

Faster, Smarter, and Safer
Coding with AI Assistants

Uberto Barbini

edited by Adaobi Obi Tulton

Early Praise for Process Over Magic

Code agents are everywhere, and they will not go away. Developers interested in trying them out should follow Uberto on his fascinating journey. This way, they will learn to use them not only for prototypes but for production-quality software.

► **Maik Schmidt**

Software Developer and Author

Uberto has done a fantastic job pulling together this guide for developers and engineering managers alike. Many of us are still scratching the surface of what AI assistance can and cannot do, and this is great first guide to bootstrap that process. Very much recommended reading for all within the development community.

► **Jayesh Shavdia**

Managing Director, Head of Markets Core Risk, Barclays

The tools we're using to create software is changing, but the principles of making systems easy to maintain are largely the same. This book provides balanced and pragmatic advice on how to use AI assistants to code faster without compromising quality.

► **Andrea Goulet**

Co-Founder, Legacy Code Rocks

The perfect "step 1" for anyone looking to understand how AI is transforming software development—enjoyable, accessible, and refreshingly honest about what works and what doesn't.

► **David Denton**

Co-Creator, [http4k](http://http4k.org)



We've left this page blank to make the page numbers the same in the electronic and paper books.

We tried just leaving it out, but then people wrote us to ask about the missing pages.

Anyway, Eddy the Gerbil wanted to say "hello."

Process Over Magic: Beyond Vibe Coding

Faster, Smarter, and Safer Coding
with AI Assistants

Uberto Barbini

The Pragmatic Bookshelf

Dallas, Texas



See our complete catalog of hands-on, practical,
and Pragmatic content for software developers:

<https://pragprog.com>

Sales, volume licensing, and support:

support@pragprog.com

Derivative works, AI training and testing,
international translations, and other rights:

rights@pragprog.com

The team that produced this book includes:

Publisher: Dave Thomas

Development Editor: Adaobi Obi Tulton

Copy Editor: L Sakhi MacMillan

Cover photo: Carina Tysvær on Unsplash

Copyright © 2026 The Pragmatic Programmers, LLC.

All rights reserved. No part of this publication may be reproduced by any means, nor may any derivative works be made from this publication, nor may this content be used to train or test an artificial intelligence system, without the prior consent of the publisher.

When we are aware that a term used in this book is claimed as a trademark, the designation is printed with an initial capital letter or in all capitals.

The Pragmatic Starter Kit, The Pragmatic Programmer, Pragmatic Programming, Pragmatic Bookshelf, PragProg, and the linking g device are trademarks of The Pragmatic Programmers, LLC.

Every precaution was taken in the preparation of this book. However, the publisher assumes no responsibility for errors or omissions or for damages that may result from the use of information (including program listings) contained herein.

ISBN-13: 979-8-88865-200-8

Encoded using recycled binary digits.

Book version: P1.0—June, 2026

Contents

	Acknowledgments	vii
	Foreword	ix
	Preface	xi
1.	Going with the Vibe	1
	Vibe Coding a Game	1
	Improving Your Prompts	6
	Writing a More Complex Game with AI	9
	What Went Wrong	13
	Using Specifications to Drive Assistants	17
	What We Covered	21
2.	Building Securely and Predictably	23
	Keep the Edge	23
	One Prompt, One Commit	25
	An AI with Good Habits	27
	Building an Application Following Rules	33
	Results and Reflections	44
	What We Covered	45
3.	Learning by Coding	47
	The Inception	47
	Shaping the Idea	49
	Walking Skeleton	52
	Modeling the Domain	55

	Developing the Search UI	57
	Show PDFs	60
	What We Covered	63
4.	Working on a Large Codebase	65
	Managing Technical Debt	65
	Understanding Legacy Code	66
	Enhancing Big Code Projects with AI	67
	Is It All Worth It?	86
	What We Covered	87
5.	Collaborating Inside an AI-Powered Team	91
	Adopting AI	91
	Team Organization	93
	Security Risks	96
	What We Covered	99
A1.	Demystifying Large Language Models	101
	Attention Is All You Need	101
	What Are Large Language Models?	102
	Vector Databases	106
	Considerations	106
	Bibliography	109

Acknowledgments

This book is dedicated to my daughter, Marina. She is growing up in a world where AI is already part of everyday life, not a novelty for her. I hope this book keeps an optimistic view of the future, and of our role in it.

My deepest thanks go to my wife, Ayumi, for her patience, support, and quiet strength while I worked on my second book. Writing one book already asks for a lot of time and attention. Writing a second one tests everyone's goodwill. I could not have done this without her support.

I am also very grateful to my editor, Adaobi Obi Tulton, for turning my stream of thoughts into clear, readable English and for reminding me to focus on what helps readers, not just what I feel like rambling about.

Finally, a big thank you to all the technical reviewers and readers who enthusiastically gave feedback, challenged assumptions, and helped improve this book. Listed in strictly alphabetical order: Paul Allton, Giovanni Asproni, Alberto Brandolini, David Denton, Andrea Goulet, Asad Manji, Ennio Masi, Duncan McGregor, Luca Minudel, Andrea Paramithiotti, Nat Pryce, Giordano Scalzo, Maik Schmidt, Jayesh Shavdia, Fraser Tooth, Matteo Vaccari, Eoin Woods, and Ayuni Zain. Your comments made the book more precise, more honest, and much better than it would have been otherwise.

Last but not least, a special thanks to Ted Neward for kindly agreeing to write the foreword.

Foreword

No six words have a more chilling effect than “Automation is coming for your job.” We just never thought it could come for us.

To some, this is karma, or cosmic chickens come home to roost. Over the last fifty years, so much of a software developer’s—particularly if you were a consultant like me—life has been centered around meeting with “subject matter experts” in a particular field, with the expressed (or barely concealed) intent to harvest the knowledge out of those folks’ brains, codify it, and deploy computers to do the job those SMEs were doing. Always, as developers, we justified it to others (and ourselves) as the price of innovation and efficiency, saying, “Well, this just enables the company to work more effectively” and “This is just the price of innovation” or “The resulting productivity gain will let you do so much more than you used to be able to.” Of course, we never stuck around to see the results of our handiwork; assuming the newly created software system actually made it to release, we were often already on to the next phase or the next project.

Then we discovered that software could do more than track inventory, weld doors to a car frame, or offer a storefront online; we discovered that software engines powered by large language models could read and/or write legal contracts, homework essays, and ... shudder ... code.

Suddenly, we as a community find ourselves on the receiving end of the treatment we handed out to others, and—not surprisingly—we don’t like it. Oh, sure, there’s a whole host of people who are excited about the power and capability that the latest round of what we so casually call “AI” offers us. “Prompt engineering” and “vibe coding” became new phrases overnight. But, if we’re really honest with ourselves, this new wave of automation strikes at a dark fear buried deep in our collective psyche: “What if they don’t need me anymore?”

All of this was top-of-mind when I sat down to read the early manuscript of this book you hold in your hands. I’m happy to report that it makes clear not

only that you're still needed but that this "AI" stuff is actually useful for us. Not "us, humanity"; I mean "us, developers."

Uberto's book makes clear something that I have believed since the first stirrings of code-generating engines powered and interfaced by LLMs came online: that these "software development assistants" provide to us the same degree of utility and benefit as the first solid-state calculators did to students of mathematics in the 1970s. To simply own one of them does not make you a whiz at mathematics. A calculator is a passive tool with no agency or insight of its own. However, in the hands of someone who understands math, who knows what formula to apply in a particular situation, and knows which values go where in that formula, that same calculator becomes an incredible source of power and productivity. Those calculators can do nothing if you take away the human punching in the numbers and operations. Conversely, humans who have learned the foundation of math are not crippled if you take their calculator away—they are simply vastly faster at their work with them. And just as the calculator gave rise to the spreadsheet, which in turn gave rise to the personal computer and later the mobile device, the current generation of tools will quickly give way to something we cannot foresee yet as we gain more experience with these tools and what's missing from them.

In the meantime, Uberto is here to show you how to get started with these tools and leverage them to make you vastly faster at your work. Enjoy.

Ted Neward

Redmond, WA, May 2025

Preface

We're living in a truly exciting moment. For the first time in history, software can answer questions, write code, and hold conversations that feel almost human. After chatting with one of them for a while, it's easy to feel like they are reasoning or understanding in a human way.

And yet, at heart, these new AI models are just ultra-sophisticated auto-complete machines. This isn't to downplay what they can do. It's simply to point out that, in a literal sense, they work by predicting what will be the next word, based on the text that came before.

They don't plan the whole answer in advance. Each new word is chosen in response to the evolving text, which means the output is built as it goes, like laying bricks one by one. Over millions of training examples, they've learned patterns about how words tend to follow one another, and they use those patterns to generate text that sounds natural and coherent.

They started out generating text but became popular in chat form. We type to them, they reply, and the whole conversation becomes the context for the next answer.

Even so, the results can be impressive. Massive training data and the ability to build "chains of thought" reasoning on the fly make them feel intelligent, but a deep analysis shows that their reasoning is a mirage.¹

They are far from perfect. They can write great code and terrible code, sometimes both in the same session. And just like when working with humans, how clearly we communicate makes a huge difference in what they produce. Clarity, structure, and readable prompts lead to better output.

What excites me most about their arrival is how much easier and faster it has become to apply good engineering practices. Techniques like TDD, modular design, and merciless refactoring were once dismissed as too slow for

1. <https://arxiv.org/pdf/2508.01191>

“real development.” Now, with AI, they’ve become often the safest and fastest way to build and evolve software.

Why This Book

I wrote this book with the goal of helping developers become more productive with AI assistants, a broad term that covers a growing ecosystem of tools powered by large language models (LLMs). This book is not about artificial general intelligence (AGI) or philosophical debates. We’ll be pragmatic here.

With the right guidance from us, LLMs can produce clean, maintainable code that fits into long-term projects. They’re not replacing developers anytime soon, but they can dramatically speed up the work. That makes them a tool every developer should learn to use well.

Alas, there’s no magic prompt shortcut. Getting good results isn’t about clever one-liners. It’s about following a solid, iterative process, step by step, just like we would when debugging a tricky issue or refactoring a tangled module. That idea—process over magic—will come up again and again throughout this book. And by the end, I hope its value will be clear to you, not just in theory but through practice.

Everything I share in this book comes from real-world experience—places where LLMs saved hours of work and others where they completely failed. You’ll learn how to get the most out of LLMs, avoid common traps, and integrate them into your workflow without losing control over your code.

These techniques work for me and my clients. Your situation might be different, so take what fits and adapt the rest. If you have ideas or improvements, feel free to reach out. I’d love to hear them.

And finally, we’re still just getting started. AI-assisted coding is in its early days. Best practices are still evolving. The goal of this book isn’t to lock in fixed rules but to give you the mindset and tools to keep learning as the field grows.

Who This Book Is For

This book is written with professional programmers in mind. You don’t need to know the exact language used in the examples, but you should be comfortable reading code and following the logic behind it. Your level of seniority doesn’t matter much—these tools are still new territory for everyone.

If you’re a technical leader or engineering manager, this book can help you figure out how to bring AI tools into your team responsibly and effectively.

You'll learn ways to keep code quality strong, design workflows that include AI, and set realistic expectations for productivity and risk.

Finally, it can also be useful if you've done some programming in the past and want to pick it up again. Maybe you're a technical manager with an idea you'd like to turn into real software. With a bit of time and effort, AI assistants can help you get started faster and take care of a lot of the technical nitty-gritty along the way.

How This Book Is Structured

[Chapter 1, Going with the Vibe, on page 1](#), walks you through two vibe coding sessions. You'll see the good, the bad, and the ugly sides of this way of working with AI. It's fast, it's fun, but it's also full of hidden traps.

[Chapter 2, Building Securely and Predictably, on page 23](#), shows how to limit the risks of vibe coding without losing the speed and flow that make it so productive. You'll learn how to set up rules—for the AI and for yourself—to keep the generative process under control. In a hands-on session, we'll build a command-line tool in Python that can index and semantically search documents using a vector database.

[Chapter 3, Learning by Coding, on page 47](#), focuses on adopting new technologies and starting new projects. We'll see how you can use AI assistants to slash the time it takes to explore and integrate new tools or frameworks. In the coding session, we'll build a nice UI for the previous CLI tool using Elixir and Phoenix, learning the technology as we go.

[Chapter 4, Working on a Large Codebase, on page 65](#), dives into working with large repositories and legacy projects. You'll learn how to guide an LLM through messy, complex code without losing track of what it's doing. AI won't magically clean up a bad codebase, but it can handle the boring, repetitive work and make the whole process a lot less painful.

Finally, [Chapter 5, Collaborating Inside an AI-Powered Team, on page 91](#), is about teams and company environments. What does it mean to work in an AI-powered team? What changes? What stays the same? We'll talk about how teams need to adjust to really benefit from these tools, and also about the new risks that come from giving AI a seat at the table.

An appendix then follows where we peek under the hood of LLMs. You'll get a simple, practical explanation of how their interference and attention mechanisms work, so you can understand how to better guide and control them. Some readers may prefer to start there to build a bit of background

before reading the rest of the book. Others may want to jump straight to the more hands-on parts and come back to the appendix only when they feel the need.

At the end of the book, just before the appendix, you'll find a one-page manual. It gathers the most important techniques from the book onto a single page you can print and keep handy as a quick reference.

Screencast Sessions

A big part of this book comes from real coding sessions with AI assistants. In those sessions, the AI wrote all the code, and my role is to act as your guide. At the start of each session, you'll find a link to the full screen recording, the total duration, and a link to the Git repository.

We focus especially on the tricky parts: why something worked (or didn't), what we can learn from it, and how you can improve your own process. To help you follow along, I've included timestamps (min:sec), screenshots, and diagrams wherever they make the flow clearer.

Across the sessions, we'll work with three different programming languages: Python, Elixir, and Kotlin, and different AI tools and IDE (in strictly alphabetical order): ChatGPT, Claude Code, Codex, Cursor, IntelliJ Junie, and Windsurf.

These examples show a range of environments, but the practices you'll learn are not tied to any specific language or tool. All modern AI coding assistants behave in very similar ways, so the techniques we cover here will apply no matter what stack or setup you're using.

You can find all the videos in the book YouTube channel.²

Online Resources

The field of AI-assisted software development is moving fast. To help you stay current and get the most out of this book, we've set up a collection of online resources on the book's companion site.³ There you'll find all the code mentioned in the book, updates, additional examples, links to new tools as they emerge, and a link to the Devtalk forum for the book.

2. <https://www.youtube.com/@ProcessOverMagic>

3. <https://pragprog.com/titles/ubaidev/process-over-magic-beyond-vibe-coding/>

You can also contact me here:

- Uberto's blog⁴
- Bluesky: @ramtop.bsky.social
- LinkedIn⁵

What's Next?

Nobody knows exactly where this is going. Some believe we're already on an exponential curve—that progress is accelerating every year and will only keep getting faster. Maybe that's true. It's definitely possible.

But, honestly, I don't think so. We're riding a big wave of innovation. It's fast, it's exciting, and it's already changing how we work. But like any wave, it won't keep growing forever. At some point, it will settle into something more steady—less explosive, more stable.

What matters most is that we start moving now. We don't need to know exactly what the future looks like to begin adapting. Developers and teams who learn to work with these tools today will be in a far better position tomorrow, no matter how fast or slow the landscape changes.

4. <https://medium.com/@ramtop>

5. <https://www.linkedin.com/in/uberto/>

Going with the Vibe

So, where do we start? Since play is one of the best ways to learn, what's better than learning how to use AI assistants by building something we can play with, literally—a game.

In this chapter, we'll explore how to work with an AI assistant and let it handle most of the heavy lifting—including writing the actual code.

This is possible with what's often called AI assistant agent mode. In this setup, the LLM isn't limited to answering questions; it can take actions. It can edit source files, run tests, write documentation, and more. Other settings also let you control exactly what the assistant is allowed to do. For example, you might allow it to edit files and run tests, but block anything that deletes data or writes to protected folders.

Working this way—without worrying too much about the code itself and simply trusting the AI to produce results—is commonly referred to as vibe coding. The idea is that you don't need to read the code or follow good engineering practices; you just go with the flow and trust the vibe.

Let's start by throwing a real challenge to the assistant: building a small but perfectly playable game, step by step.

Vibe Coding a Game

For our first experiment, we'll have the LLM create a Jewel Swap game in Python. We'll use Pygame, a handy library that makes writing games much easier. It's simple and quick to set up, and it's popular—so the LLM has likely seen plenty of similar examples during training.

Preparing the IDE

For this project, we'll use Windsurf as our AI-assisted IDE.¹ It's the one I personally prefer, but of course, there are other solid options like Claude Code,² Cursor,³ Copilot,⁴ or Cline.⁵ Things move fast, and the tools keep evolving.

Windsurf is a stand-alone IDE based on a fork of Visual Studio Code. A plugin is also available if you're using IntelliJ or PyCharm.

In the end, the choice doesn't matter too much. These tools work in roughly the same way, and even if some details will change over time, the general approach we'll show here should still apply.

We'll be using its Cascade chat interface with agent capabilities, which means we don't have to run every command ourselves; the agent will handle that part. It's productive and convenient because it lets us focus on the code instead of constantly switching to the terminal.

The Jewel Swap Game

Let's start fresh with a brand-new project and ask the model to write the game from scratch. What you'll see next is the actual chat I had with the assistant—no edits, just my prompts and its responses. The full code is in the book's repository, and if you'd like to see how it unfolded in real time, you can also watch the screencast.⁶

Feel free to copy the prompts and try them out with whatever tool or model you're most comfortable with. Your version of the game will probably look a little different, but you'll probably run into the same kinds of bumps and breakthroughs along the way.

All we need now is a clean project folder, Python and Pygame installed (just ask the assistant if you're not sure how), and the assistant running in agent mode. Let's see how it handles the challenge.

> Can you write a game like Jewel Shuffle in python using pygame for my daughter? Please make it cute and playable with mouse!


Checked command status

-
1. <https://codeium.com/windsurf>
 2. <https://www.anthropic.com/claude-code>
 3. <https://www.cursor.sh>
 4. <https://github.com/features/copilot>
 5. <https://cline.dev>
 6. <https://youtu.be/7DV4MvSR1v0>

I'll create a cute Jewel Shuffle game for your daughter using Pygame with mouse controls. Let me first check what you already have in your project.

...

As you can see from the following screenshot, there's quite a bit of verbosity here. The assistant is doing a lot: checking the files, editing them, trying to run the game, and when things break, attempting to fix the issues on its own.



After a few hiccups with imports (around the 01:50 mark in the screencast), everything came together. The game was running, and the assistant proudly let us know it was ready:

Checked command status

I've created a cute Jewel Shuffle game for your daughter!

The game is now running and should be playable with the mouse.


Game Features

Cute Design: The game features colorful jewels with different shapes (hearts, stars, diamonds, flowers, moons, suns, and circles) on a soft pink background.

Mouse Controls: Click on a jewel to select it, then click on an adjacent jewel to swap them.

...

Less than two minutes, and the assistant finished the game! All the code is packed into a single file, but it runs. And while it's not going to win any design awards, the pastel-colored tiles look pretty nice, as you can see in the following image.



It's genuinely amazing the first time you see an LLM generate working code just from your words. You type a request in plain English—no syntax, no structure, just a casual description of what you want—and the assistant not only understands it but starts writing real code that does the thing. It feels like unlocking a secret door.

Suddenly, the wall between idea and implementation gets thinner. The model picks up on context, catches the nuances, and sometimes even guesses what you meant better than you could've explained it.

The first time it happens, it's hard not to sit back and think Wait ... did that just work?

Unfortunately, reality creeps in fast. Not all that glitters is gold. The code might look good at first glance, and the assistant's confident explanations make it sound flawless, but as soon as you try to run it, things start to break. The generated code is far from perfect, and it often doesn't even match its own optimistic description.

As soon as we try to play this jewel game, a problem pops up: some tiles don't fall correctly, and they get drawn on top of each other.

No reason to worry; we can just ask the assistant to fix it!

> there is a problem on some tile drawing and falling correctly

I'll fix the drawing and falling issues in the game. Let me examine the code to identify the problems:

Analyzed

```

main.py:(1-200)
Let me check the rest of the code to identify the issues with drawing and
falling:

Analyzed

main.py:(201-400)
Let me examine the rest of the code to fully understand the drawing and
falling issues:

Analyzed

main.py:(401-600)
After reviewing the code, I've identified a few issues with the drawing and
falling mechanics. Let me fix these problems:

The main issue with falling is in the remove_matches method, where matched
jewels are only removed when their alpha value is low enough, but we're not
waiting for the fade animation to complete before applying gravity.
There's also an issue with the drawing of jewels - some shapes might not be
rendering correctly.
Let me make the necessary changes:

...

```

Okay, at this point the assistant seems to have a decent grasp of what's going wrong and why. Notice how it analyzes the code in chunks—roughly 200 lines at a time—and pinpoints where the issues are.

It updates the `remove_matches` and `draw_jewels` functions, and (around the 05:30 mark) it proudly announces that everything is fixed.

Only ... it's not. How tiles are swapped still has problems: if the swap doesn't produce a merge, the tiles should go back to the original position, but they don't.

Time to get more specific with our instructions.

Also, take note: writing the game from scratch took just 1.5 minutes. Fixing it? Over 4 minutes, and it still didn't get it right. Debugging is clearly the harder part, even for an LLM.

> when I swap 2 jewels the second doesn't go where the first was

```

I'll fix the issue with the jewel swapping. The problem is in the
swap_jewels method where the jewels aren't correctly swapping positions.
Let me fix that:

...

```

I've fixed the issue with the jewel swapping. Now when you swap two jewels, they should correctly move to each other's positions. The problem was in the `swap_jewels` method where the position updates weren't being handled properly.


Here's what I changed:

```
Instead of using Python's tuple unpacking to swap the grid references, I now:
Store the original positions of both jewels
Explicitly update the grid references
Set the target positions for each jewel correctly
This ensures that each jewel properly animates to its new position, and
the grid references are updated correctly.
```

The game should now work properly with jewels swapping correctly.
The animations should be smoother, and the falling mechanics should also work better with the other fixes we made earlier.

Is there anything else about the game that needs adjustment?

Alright, the swap bug is finally fixed! The game still isn't complete and minor issues remain, but it's now playable, and that's already a solid win (see the following image). It's only seven minutes since we started!



And honestly, it's still pretty amazing that we were able to build the core of a playable game in just a few minutes. Doing the same manually would've taken me several hours, even considering I already know Python and Pygame.

Improving Your Prompts

I've run the preceding little experiment a few times now, using different models and AI assistants. The code and graphics varied each time, but the end results were surprisingly similar: they all reached something playable pretty quickly and then hit a wall when it came to polishing and refining the details.

But don't just take my word for it. Try doing the same exercise yourself, using different models and assistants. It's a great way to get a feel for their strengths

and limitations. It's also a solid exercise for learning how to express our requirements more clearly, so we can guide the assistant toward better results.

Being able to write clear instructions that help AI assistants push through difficult tasks is a skill in itself. It takes practice to get right, and it's often referred to—perhaps a bit pompously—as prompt engineering.

This raises a bigger question: are LLMs fundamentally unable to handle tasks beyond a certain level of complexity without human guidance, or is this just a temporary limitation that future models will overcome?

Over the past two years, progress has been steady. LLMs are definitely getting better at writing cleaner code and fixing issues more quickly. But the core experience hasn't changed that much since my very first attempt, back when the first version of Cursor came out—the first real chat-based coding assistant.

One common strategy for tackling harder problems is called chain-of-thought reasoning. In this mode, the LLM uses your prompt to spin out a much longer, step-by-step explanation of how it plans to solve the problem. It then takes that self-generated reasoning and uses it as a new prompt to write the code. Sometimes it even repeats this process recursively, adding extra rounds of “thinking” to check and improve the result.

The problem is that while the reasoning often looks solid, the final outcome is far from guaranteed. Even with a long chain of thought, the assistant may fail to solve the problem. It also takes longer to reach an answer, and sometimes the model gets stuck in an endless loop, unable to produce a result it considers “good enough.”

For that reason, I usually reserve this deep-thinking mode for the hardest problems, and I let it run in the background while I do something else. For everyday tasks, it's usually faster and easier to just check the output myself and correct mistakes on the fly.

The trend with new models has been consistent: every few months a highly hyped release arrives, boasting improvements in areas where LLMs are already strong—supporting more languages, frameworks, and algorithms. But the same core issues remain. They still produce code that doesn't compile because of calls to nonexistent methods. They still follow instructions too literally without solving the real problem. And they still make frustrating mistakes no human developer would make, like changing test assertions just to get them to pass. These flaws were there in the early days, and they're still with us today—even in the most advanced models.

The good news is that we're not powerless. By changing how we interact with these systems, we can greatly reduce the frequency of those mistakes. When prompting an LLM, it's better to think less like you're talking to a person and more like you're designing an input that maximizes the chance of getting the output you want. The mindset shift is from conversation to probability shaping.

- *First, guide it—don't try to convince it:* The model doesn't “understand” you, it matches patterns. Your wording, structure, and constraints are what push it into the distribution of answers you're aiming for.
- *Second, start with the shape of the result:* Before you write the prompt, picture exactly what you want back—a list, a code snippet, an essay—and then design your prompt to funnel the model there. Adding constraints like “four numbered bullet points and one code example” makes the target clearer.
- *Third, use structure and roles:* LLMs respond well to formats like Markdown lists, JSON schemas, or role instructions (“you are a security expert reviewing code”). Structure gives the model something to lock onto.
- *Fourth, prime with context:* Examples, audience cues, or exclusions (“skip theory, focus on steps”) help narrow its path.
- *Finally, iterate like an engineer:* Test, tweak, compare, and reuse prompt snippets that work. Prompting is less artful conversation and more structured design.

Does It Pay to Be Polite?

As you may have noticed, I usually keep a polite tone with AI assistants. I use please, thank you, could you, and so on. Is it a good idea?

For me, yes. It helps me stay relaxed and focused. When I switch to short, sharp commands, I end up in a kind of “dictator mode,” and that hurts my attention and my coding skills.

Interestingly, there's also some evidence that how you phrase a request affects the quality of the response. A 2024 paper “Should We Respect LLMs? A Cross-Lingual Study on the Influence of Prompt Politeness on LLM Performance” (<https://arxiv.org/abs/2402.14531>) tested prompts from very polite to openly rude across multiple tasks in English, Chinese, and Japanese.

Impolite prompts led to worse results: more errors, stronger bias, and missing information. In some cases, the models even refused to answer when the prompt was too disrespectful.

Writing a More Complex Game with AI

Andrej Karpathy (former Tesla AI director and founding member of OpenAI) coined the term *vibe coding* in a tweet on February 3, 2025, to describe a kind of lazy approach where you stop caring about what's happening under the hood and just rely on the AI, accepting every suggestion to get the task done as quickly as possible.

He clearly defined it as not really coding: “I just see stuff, say stuff, run stuff, and copy-paste stuff, and it mostly works.”

Since then, the term has spread like wildfire. For many, *vibe coding* has become the new gospel: they preach that with AI, there's no need to understand—or even look at—the code anymore. That approach can lead to all sorts of problems, as we're going to see.

Let's try something a bit more challenging.

So far, we've seen how the assistant can handle a simple puzzle game, but what happens when we ask for something slightly more complex, something that requires more logic, physics, and interaction?

Let's ask it to write a basic flight simulator. Nothing fancy, just a simple 3D simulation where we can control an aircraft, and maybe deal with gravity and some basic flight mechanics. This will push the assistant beyond drawing tiles on a grid and into territory where structure, reasoning, and correctness matter a lot more.

Let's see how far we can get. Watch the video.⁷

```
> Please write the code for a flight simulator of jet fighter on a mountain
> ground. The goal of the game is to reach the goal avoiding the radar
> staying as low as possible. Use pygame and any library you think can help.
```

```
I'll create a flight simulator for a jet fighter navigating through
mountains to reach a goal while avoiding radar detection. Let's use pygame
for the graphics and physics simulation. I'll also use NumPy for
mathematical operations and terrain generation.
```

```
...
```

```
README
```

```
Jet Fighter Flight Simulator
```

```
A 3D flight simulator where you control a jet fighter navigating through
mountainous terrain. Your mission is to reach the goal while avoiding radar
detection by flying as low as possible.
```

7. <https://youtu.be/roO8VF83j4>

Features


- 3D terrain with mountains generated using OpenSimplex noise
- Realistic flight physics with pitch, roll, and yaw controls
- Radar detection system that is more likely to spot you at higher altitudes
- Heads-up display (HUD) with altitude, speed, attitude indicator, and more
- Goal-seeking navigation system

...

I've created a 3D flight simulator game where you control a jet fighter navigating through mountainous terrain. The game is now running, and you can see it's working though it seems the initial position might be causing the jet to crash into terrain immediately - this can be adjusted by changing the starting position in the main.py file if needed).

It finished in less than five minutes, and after reading that description (around the 04:30 mark), we are super-excited. The LLM started generating files with names like Camera, HUD, Radar, and more, each filled with what looked like a lot of promising code.

But trying to run the game was a letdown. It crashed immediately, and when we finally got it to start, the “landscape” was just a scattered bunch of dots, as you can see in the following image.



Still, no need to panic. The first version of the jewel game wasn't perfect either. So let's do what we did back then: ask the AI to fix it and see what happens.

> can you improve the terrains? it's just dots now

I'll improve the terrain rendering to make it more visually appealing

instead of just dots. Let's modify the `terrain.py` file to render triangles or polygons for a more solid terrain representation.

...

The game is now running with the improved terrain. You can see solid mountains and valleys instead of just dots, which should make navigation more intuitive and visually appealing.

Let me know if you'd like any further improvements to the terrain or other aspects of the game!

It took a while, but after a few tries (around the 09:10 mark) everything compiles and we were ready to try running it again.

Well ... it doesn't crash immediately this time, which is progress. But the controls don't behave like anything you'd expect from a flight simulator. Even worse, the terrain is now a brown blob—definitely not the mountains and valleys we were promised.

> now I can only see a brown blob

...

The terrain should now appear much more visually interesting with:

A proper sky gradient from deep blue to light blue at the horizon
 Varied terrain colors based on elevation (grass, dirt, rock, snow)
 More natural-looking mountains with randomized shapes and sizes
 Better camera positioning to see more of the terrain


After running for about three minutes, the assistant confidently told us that everything was fixed and ready to go (12:00).

You're probably already skeptical, and rightfully so. As you can see in the next screenshot, what we had now was just a slightly different brown blob as terrain. The controls still don't behave like a flight simulator, the graphics are rough, and on top of that, everything runs painfully slow. The game is still completely unplayable, as you can see in the image on the next page.

This was a very different experience from the jewel game. Not only were we nowhere near our goal, but with each iteration, things didn't seem to get better at all. In fact, the code just kept growing—hundreds, even thousands of lines we don't understand and we can't trust.

At this point there's no point in continuing. Starting over from scratch would've probably been easier than trying to untangle that mess.

And just to be clear, this screencast isn't a worst-case scenario. I ran this multiple times, and what you're seeing here is actually one of the better results.



You might think it's unfair to ask the assistant for a flight simulator. But here's the thing: it never pushes back. Not once does it say, "This might be too ambitious" or suggest we lower our expectations.

It didn't pause to check if things are working before piling on more code. Instead, it cheerfully asked if we wanted to add new features, as if everything were running smoothly.

And that's part of the problem: if you're not already an experienced developer, it's hard to know what's simple and what's very complex in software development. The assistant's tone gives no hint. It treats "add shadows" or "simulate aerodynamics" the same way it treats "change background color," and that can be misleading. Without that intuition, you might not realize you're heading straight into a wall until it's too late.

Once you experience this kind of "dark vibe" behavior—the assistant confidently marching forward while everything's broken underneath—it shatters that initial rose-tinted view. And honestly, it's no surprise some people walk away thinking the whole thing is a waste of time, only good for throwaway experiments.

But that would be a mistake.

Building a simple yet engaging flight simulator is absolutely possible with an AI assistant. You just need to know how to guide it properly, and how to give it the right instructions.

That's exactly what we dive into in the next chapter. But first, let's take a moment to reflect on what just happened.

What Went Wrong

Let's take a step back and go over all the issues we encountered so far and what they reveal about working with AI assistants.

The goal here isn't just to point fingers at the model. Instead, we want to understand these pitfalls so we can build a smoother, more reproducible process that avoids them as much as possible.

Hallucinations

One of the first and most obvious problems is the assistant often tried to run code that couldn't possibly work: missing imports, nonexistent methods, calling functions that had never been defined.

This kind of issue is commonly known as a hallucination. The LLM generates code that looks right—function names sound reasonable, class structures seem familiar—but they're often completely made up. Sometimes it pulls inspiration from other languages or from training examples that don't fully match the current context.

This happens because of how LLMs work. They don't "know" facts—they predict the next most likely word based on patterns they've seen during training. When the model doesn't have a clear answer, it still tries to fill the gap with something that sounds right. In short, hallucinations aren't magic or mystery—they're a statistical side effect⁸ of a system designed to be fluent rather than certain.

To its credit, in our examples the assistant managed to fix these issues by itself. But they still eat up time and disrupt the flow, especially considering it should at least handle the easy parts.

The Drunken Intern

Sometimes, the assistant would confidently say that everything was working even when it clearly wasn't.

This is what I like to call the drunken intern problem: it's cheerful, eager to please, but stubborn and completely wrong.

One reason for this behavior is that LLMs cannot see what's happening on the screen when running the code. They have no feedback loop from the actual runtime behavior of the game.

8. <https://arxiv.org/html/2509.04664v1>

Another likely reason is pattern-matching bias. The model has seen a lot of working examples during training, so it tends to assume that similar-looking code will behave the same way. But as we know, with code even tiny changes can lead to very different outcomes.

For the same reasons, it's also very tricky to measure LLMs' capabilities. As the paper "What Does Human Evaluation Even Mean?"⁹ explains, LLMs often outperform humans in tests and benchmarks not because they "understand" more but because they exploit patterns and shortcuts in the tests themselves.

Losing Context

At times, the assistant seemed to lose track of its own work. It would forget what parameters a method had, or how a helper function was structured, even though it had just written them.

This usually comes down to context window limitations. The model can only "see" a certain number of tokens at a time. If the project grows beyond that, it starts dropping details, and you get inconsistent or broken suggestions.

The good news is that context windows are getting bigger with every new model release. This issue may fade over time, but for now it's something to actively manage when working with AI on larger codebases. We'll see this issue again when discussing [Working in Small Steps, on page 71](#).

The Loop of Death

This is one of the most frustrating traps you can fall into when working with an AI assistant, especially on bigger projects.

At some point, the model starts fixing one thing while accidentally breaking something else. Then it tries to fix that new problem and breaks the first thing again. And so on. Welcome to the loop of death.

Another painful variation is that you ask it to add feature X, but it also sneaks in some unwanted side effect Y. When you tell it to remove Y, it also removes X. And you're back to where you started.

In these cases, the assistant just keeps cycling through partial solutions, without realizing it's already tried them. It kind of hopes that if it spins the wheel enough times, something will magically work.

9. <https://arxiv.org/pdf/2502.07445>

When you hit this loop, there's only one real solution: Stop. Step back. Investigate manually. If you can understand what's going on, you can give it a much better prompt that points directly at the solution.

Without tests and a clear understanding of the code, this can turn into a huge time sink, especially if you're doing vibe coding, chasing ideas without a tight feedback loop.

It's a hard but important reminder: LLMs don't actually understand the code. They're just predicting likely patterns. And when the context gets too fragmented, their guesses start to crumble.

Random Results

It's not exactly a failure, but if you run the same prompt multiple times, you'll often get very different outputs.

That's simply how LLMs work—their generation process isn't deterministic. Combine that with a short or vague prompt, and you can end up with wildly different applications each time.

This unpredictability can be fun when you're exploring or looking for creative ideas. But for serious work, we need more control. Ideally, we want to guide the process in a way that's both reproducible and predictable.

Low-Quality Code

Even when the code works, there are often structural problems in the AI-generated code that you need to be aware of.

Generated code is often affected by tight coupling between modules, duplicated logic, unclear separation of concerns, messy naming, and so on—in short, low internal quality.

Now, you might think, “Who cares? If it works, it works.”

But it doesn't work like that. Poor-quality code is much harder for LLMs to fix or refactor later. It increases the chances of falling into traps like the loop of death we talked about earlier. And once the code becomes unreadable, it slows you down too, especially when you have to step in and fix things manually.

Paradoxically, bad code creates more problems for the LLM than it does for a human. We can often reason our way through a mess. The assistant, on the other hand, just gets confused.

Performance Traps

Even when the code runs fine, it might be hiding all sorts of inefficiencies: unnecessary or repeated function calls, poor use of data structures, ignoring language-specific features, or completely missing better algorithms and high-level optimizations.

If you're running that code in the cloud, those inefficiencies can quickly become a real money problem. Plenty of developers have been surprised by massive bills, all because they deployed LLM-generated code without proper checks or controls.

Hidden Costs

Talking about bills, another thing to watch out for is how fast costs can add up when you're building with agents and big models, especially when you're using a vibe coding approach with multiple agents, each one checking different things.

It's easy to underestimate how much is being processed behind the scenes. You might think you're just asking for a small change, but if the assistant has to load the full codebase, track multiple goals, and spin up several sub-agents, you could be looking at thousands or even millions of tokens in a single session.

Security Nightmares

Here's another hidden danger: huge security holes. LLMs can easily leave sensitive data, like API keys or credentials, right in the source code. They might generate functions that look useful but introduce common vulnerabilities. Even worse, they sometimes copy insecure practices from outdated or bad examples in their training data.

Known attacks have even exploited this behavior: malicious packages uploaded with names that match common hallucinations, specifically designed to be picked up by LLM-generated code.

And since the code technically “works,” these risks often go unnoticed, unless a human takes the time to review it properly.

Missing or Misleading Observability

Observability is often treated as an afterthought in AI-generated code, if it's included at all.

Many times, LLMs skip logging entirely. Other times, they insert log statements that look helpful but don't reflect what's going on. And sometimes they're flat-out wrong, saying a step succeeded when it didn't or logging the wrong variables entirely.

Good observability is a nonnegotiable part of building reliable systems. Make sure to check if your code works correctly.

The Mindset Trap

Maybe the most dangerous issue isn't in the code at all but in how we think.

An interesting study on GitHub repositories and tools like Copilot¹⁰ found a clear increase in copy-pasted code and a drop in code quality even in production projects.

The study highlights some specific risks AI assistants introduce:

- More redundancy and duplicated logic.
- Developers blindly accepting suggestions without fully understanding them.
- A decline in critical thinking and hands-on problem-solving skills.

This mindset shift—letting go of control and curiosity—can have long-term consequences on the whole product.

Using Specifications to Drive Assistants

One way to reduce some of these problems is to use Spec-Driven Development. Instead of relying only on vibe and letting the assistant run wild, we slow things down a little and add some planning. It's not as fast or carefree as pure vibe coding, but it gives more reliable results.

We start by writing a clear, detailed specification of what we want the app to do. Then, instead of asking the assistant to build everything at once, we break the work into small, manageable steps.

For each step, we let the assistant generate the code freely. Then we run the app and check that it behaves as expected. If it works, we move on. If it doesn't, we ask for a fix until we're satisfied. We also ask the assistant to keep the documentation up-to-date as we iterate.

10. <https://arxiv.org/html/2412.06603v2>

This method doesn't require reading the code, and it's very fast. It doesn't even require a developer to get started. Anyone with a clear idea can build and test a minimal version of their product quickly and cheaply. It's no surprise that it's so popular.

The results are more consistent, but it has the same risks of pure vibe coding: even if the application is working as we want, there could be security leaks or other problems hidden in the code.


Spec-Coding a Little Platform Game

At the start of the session—watch the video¹¹—we gave the brief and asked the AI to expand it with all the technical requirements. This example is small, but real-world specifications can easily become very long and detailed.

The assistant immediately wrote a plan and started moving through the steps one by one. This is the so-called Planning Mode, where the model runs a kind of internal conversation: first it drafts a plan, then it follows it, and finally it checks whether the result matches the expectations—revising the code when it doesn't.

While grinding through the steps, it generated a large amount of code. We didn't have to do anything. But around the 1:23 mark, I noticed that the level map was completely wrong for a vertical platformer, so I stopped the run and asked it to fix the design.

Once it adds lives management and a simple menu, we have a complete game; not a very exciting one—the gameplay is almost nonexistent and the graphics are just placeholder squares—but all in all, not a bad result for about 20 minutes of work (see the following image).



11. <https://youtu.be/atpjMv336RE>

Command-Line Assistants

If you watched the video, you saw that I used Codex,¹² OpenAI’s command-line assistant, for this exercise. Unlike Cursor or other IDE-based tools, these assistants (Claude Code, OpenCode, and similar ones) run in the terminal and behave more like a specialized shell than an editor.

For projects where we don’t care much about inspecting the code directly, this makes a lot of sense. The terminal view makes it easy to see what the assistant is doing and to give instructions without distractions.

You can also see indicators showing the current state of the LLM, memory usage, and so on, as shown in the following screenshot.

```

ubertobarbini@oryx-laptop:~/develop/python/panda-platform$ codex
> _ OpenAI Codex (v0.58.0)
model:      gpt-5.1-codex  /model to change
directory:  ~/develop/python/panda-platform

To get started, describe a task or try one of these commands:
/init - create an AGENTS.md file with instructions for Codex
/status - show current session configuration
/approvals - choose what Codex can do without approval
/model - choose what model and reasoning effort to use
/review - review any changes and find issues

> I want to create a platform game, very simple, like retro 8bit games. The protagonist will be a Panda and it can walk, jump (not very high) and climb bamboo. Can you write first a detailed specification for this game so I can review it
100% context left

```

It’s also possible to run multiple AI agents in different terminals at the same time. But apart from the risk of burning through your token quota very quickly, you also risk conflicts and merge issues if they work on overlapping areas.

Is Code Now Irrelevant?

Some people, like Sean Grove from OpenAI, argue that just as modern languages are abstractions over assembly, specifications could become abstractions over code.¹³ According to this view, we could treat code as a temporary artifact (like compiled binaries) and use the specification as the real source of truth.

With all due respect, I don’t agree that “code is a lossy projection of specifications.” A lot of knowledge doesn’t live in the spec. It appears during the iterative process of writing, refining, and understanding code.

We can safely ignore compiled binaries because we can always rebuild the same machine code from the same source and the same compiler version.

12. <https://openai.com/codex/>

13. <https://www.youtube.com/watch?v=8rABwKRsec4>


But this isn't true with LLMs. They are nondeterministic by nature. They'll never produce the same result twice.

As Eoin Woods put it when reviewing this book: “Layers of formally defined languages and reification between them is *totally* different to a nondeterministic translation from natural language to formal language.”

Running the Same Spec Twice

To push this idea further, I asked the assistant to update the specification with everything we had discussed. Then I copied the spec into a new directory, opened a fresh session with the same assistant, and asked it to build the game again from scratch. Watch the video.¹⁴

Interestingly, the level map is still wrong for a vertical platformer, but aside from that, the run produces another small game—a different small game (see the following image). It's better in some aspects (I like the little panda eyes, and the scrolling feels smoother in this version) but worse in others (the enemies don't really work).



Both runs produced a reasonable interpretation of the brief and the specification, but they did it differently. It's difficult to imagine a natural-language specification precise enough to avoid these kinds of variations. And even if it were possible, how large would that document be?

Are we sure replacing a well-organized codebase with hundreds or thousands of pages of Markdown documents would be an improvement? Which one would you rather browse to find a bug?

Finally, as impressive as the results are, they're still far from a ready-to-play game. Even if they are technically “correct,” playing them for a few seconds makes you cringe at all the missing little details that make a game fun and engaging.

14. <https://youtu.be/78-302HxG5g>

Maybe one day assistants will turn specifications into production-quality code as reliably as humans do. But that will require new breakthroughs. LLMs are not there yet.

Vibe Coding and Spec-Driven Development: The Good Parts

Having said all that, let's close this chapter on a more positive note.

Vibe coding is exciting. Sometimes, in just a few minutes, the assistant gives you a complete, working application. If you're lucky, it just flows. That first time is a real wow moment, and it's no surprise many people get hooked.

It's especially great for writing proofs of concept. What used to take days can now be done in a couple of hours. That makes it perfect for quickly exploring a new technology, testing out an API, or experimenting with an algorithm you've never tried before.

In general, it's totally fine to use vibe coding for things you don't care too deeply about—internal tools, weekend projects, one-off utilities. You're not shipping these to external users, and it's okay if they're a bit slow, messy, or unsafe.

It also works surprisingly well for small, common applications—things like a mortgage calculator, a to-do list, or a basic CRUD interface. These patterns show up so often in training data that LLMs tend to get them mostly right out of the box.


This isn't a binary choice; it's a spectrum. On one end you have fully controlled development, on the other pure vibe coding. You don't have to commit to an extreme. Depending on the task, you might use broad, high-scope prompts or very focused requests, and you might review every line, skim through, or skip the check entirely. All of these approaches can be valid, as long as you're conscious of the trade-offs and risks that come with each.

The diagram at the top of the next page shows how the different ways of using AI assistants relate to each other along two dimensions: prompt task size, which is how large a feature you're asking for at once, and generated code review, meaning how carefully you check the code before committing it. Each style has its place, so there's no single "right way."

Still, be careful—it's surprisingly easy to drift into a complacent, sloppy mode. Keep your eyes open to what's going on in the code.

What We Covered

In this chapter, we had our first hands-on experience with an AI assistant.



- *AI can deliver impressive results:* With the Jewel and the Panda examples, we got a playable game in just a few minutes without touching the code ourselves. That was impressive, even if there were still issues here and there.
- *AI can fail miserably:* The second example, trying to build a flight simulator, was a different story. It was a more complex task, and we didn't guide the AI much on how to handle it. The result was deeply disappointing. Not only was the game not working but the assistant kept insisting that everything was fine. It even suggested adding new features, completely ignorant to the fact that the game wasn't playable at all.
- *Lack of process leads to chaos:* In these examples we completely trusted the AI, throwing ideas at it and seeing what came back. For the rest of this book, we'll take a very different approach. Instead of loose, back-and-forth conversations, we'll rely on discipline and solid engineering principles to get more predictable and reliable results.

In the next chapter, we'll start with a set of practices to help keep the risks of working with AI assistants under control without giving up (much of) the speed and convenience that make LLMs so appealing.

Building Securely and Predictably

In the last chapter, we saw that vibe coding, while fun and sometimes surprisingly effective, doesn't cut it for applications that need to run reliably and safely in production. It's great for quick experiments, internal tools, and one-off projects. But when we care about maintainability, performance, or security, we quickly hit the limits of this approach.

So what's the alternative? Should we go back to writing everything by hand? Luckily, no. Otherwise, this would be a very short book.

In this chapter, we'll explore how to guide AI assistants to produce solid, reliable code—not just prototypes, but production-quality software. To do that, we also need to grow our skills as AI supervisors, learning how to steer the assistant, review its output, and keep the process on track.

That means applying some discipline. Slowing down a bit. Making sure each step makes sense before we move on. But this doesn't take the fun away. In fact, working this way can be even more rewarding. We keep the fast feedback loop, the creativity, and the flow, and we end up with code we're happy to maintain.

Let's dig into how to make that happen.

Keep the Edge

When I was a kid, my father used to take me skiing in the Italian Alps. While he was off with his friends, I was stuck in ski school, which I absolutely hated.

But there's one thing I still remember clearly: my instructor constantly shouting at me, "Keep the edge!" At the time, I didn't get it.

I was just a kid, and all I wanted was to go fast. Straight down the slope, no turns, wind in my face—pure speed. Like flying. But those runs always end

in a crash, and crashes in skiing can be dangerous. My instructor saw it coming. He kept yelling at me to use the edge of my skis, to grip the snow and stay in control. Because once you lose the edge, the skis go wherever they want and you're just along for the (usually short) ride.

Now, working with LLMs, that phrase keeps coming back. “Keep the edge”—because it’s the same feeling. It’s tempting to go full speed: automate everything, chain agents together, let the assistant push code straight to production. But the moment we ease off and stop checking what’s going on, things can spin out quickly.

Sure, nobody breaks a bone if the code doesn’t compile, fortunately. But it’s still painful to realize that your entire project is a pile of junk after hours of “vibing” work.

Speed feels great but only if we stay in control. The edge is what keeps the ride smooth.

So ... what now?

Can we still get real value from AI assistants without risking the whole house?

Yes, absolutely.

Do you need to be a programmer to do it?

Not necessarily, but you do need to think like one. In other words, the better you understand your problem, and the steps needed to solve it, the better your results will be. AI tools can be incredibly powerful, but they’re not reliable. You still need to lead the way. And that means giving them clear, detailed guidance if you want consistent, high-quality output.

To move beyond the limitations of vibe coding, we need to shift our mindset. Instead of throwing vague prompts at the assistant and hoping for magic, we have to provide clear, structured instructions and double-check everything it gives us. In other words, the trick to making LLMs work for coding isn’t finding some magical “perfect prompt;” it’s about setting up a solid process and sticking to it.

But even before we start “training the assistant,” we need to do something even more important: we need to define a set of practices for ourselves—because how we work with the AI matters as much as, if not more than, what we ask it to do.

One Prompt, One Commit

The most important rule to follow when working with generative AI is what I call One Prompt, One Commit. This means that every time you prompt your AI assistant, you should ask for something specific, small, and self-contained. Once it's completed and verified, you commit it to version control before doing anything else.

Mastering Git

If you want to get the most out of working with AI assistants, it's worth knowing Git well—especially how to move between commits and branches without mistakes. Each commit should be a safe checkpoint, completely atomic—focused on a single, complete change. That way, if the assistant breaks something, you can roll back cleanly without losing unrelated work.

Since working with AI often means exploring parallel ideas and backtracking, this is a good time to refresh your Git-fu. Using branches lets you isolate experiments and work on them in parallel.

You don't need every Git command, but being confident with `diff`, `branch`, and `worktree` gives you the safety net to experiment freely while keeping control.

In particular, `git worktree` is a little-known gem that lets you (and the AI) work on multiple branches at the same time—without constantly checking out or stashing changes.

It allows you to check out several branches, each in its own subfolder, making it ideal for experimenting: you can compare results quickly, merge the good parts, and keep your main branch clean.

This rule builds on four implicit assumptions. These are good practices in any development workflow, but they become absolutely essential when you're working with AI-assisted code:

- *Only commit working code:* You should only commit code that builds, runs, and passes all relevant tests. That means writing tests—both unit and end-to-end—for every new piece of functionality. If the code doesn't pass the tests, don't move on. Fix it first. Conversations with the AI are fine at this stage, but they must stay focused on fixing the issue, not on adding features or expanding scope.
- *Always review before committing:* Before each commit, review all changes. Make sure the code does exactly what it should and that nothing unexpected slipped in. If there's a critical issue, stop and fix it. If it's a minor design or readability tweak, note it for the next prompt and commit cycle. Resolve one issue at a time. In other words, never commit code you're not

sure about. If unsure, you can always ask for an explanation from the AI assistant itself.

- *Focused prompts:* Each prompt and each commit should focus on a single concern. Either you're adding a new feature or you're refactoring—not both. Don't mix goals. Mixed intentions lead to messy diffs, confusion, and painful rollbacks.
- *When stuck, roll back:* If you get stuck and the assistant can't seem to fix the problem, don't waste time fighting it. Roll back to the last clean commit and try again with a smaller, more focused prompt. This helps you avoid falling down a rabbit hole of broken logic and unclear state.

You might be thinking, “Sure, nothing surprising here. This is just solid development hygiene.” And you'd be right.

But the key difference is that when working with LLMs, these practices aren't just helpful, they're essential to keep your project from going off the rails:

- Small commits make it easy to roll back at any point without fear of wasting time.
- AI assistants behave more reliably and less randomly when given focused, detailed prompts.
- When bugs appear, it's easier to trace them, since each change set is small and clear.
- You avoid hidden side effects that are harder to catch in large, complex diffs.
- Small focused commits make it easier to collaborate with other developers or even just remember, the next day, what you were doing.

These constraints might feel heavy at first, but they bring much-needed structure to your workflow. In the long run, they save time and help you consistently deliver what you need.

Do you need to edit files manually during this process? Usually not. It's not a hard rule, but I rarely touch the code directly when working on something new. That said, sometimes it's faster to fix things in the code rather than explain what to do to the AI agent. You can adapt the process to suit your style. For example, you might write the tests or the class skeletons yourself and let the LLM fill in the details.

So how tightly should you review the code before commit? Ultimately it depends on what you're doing. For production code you should check the code line by

line just like a normal PR review from someone you don't fully trust. For non-critical projects, you can be more relaxed and just skim the code for anything that looks wrong. Tests are helpful, but they can't catch everything.

But make no mistake, this isn't vibe coding. Even if the AI agent is writing all the code, the experience is very different from the loose, free-form style we used in the previous chapter. This approach demands focus. It can be mentally tiring, like keeping the edge while skiing downhill.

The speed of development is high, but it requires constant attention and judgment. In fact, many developers find they need more breaks than usual just to keep up.

On the other hand, this approach is massively productive. It's hard to put exact numbers on it, but developers using this process can routinely complete in a few hours tasks that would've taken days without AI assistance. And it's not just about speed; since we can spend more time on refactoring, testing, and cleaning things up properly, the overall quality usually ends up higher too.

An AI with Good Habits

And like with human developers, good habits make all the difference for the LLM. So how do we tell the LLM which habit to follow? Well, we ask the LLM to write it down in a special document called the ruleset.

What's the Ruleset?

All LLM-based assistants allow you to define a persistent prompt that gets automatically appended to every request. This is often written in a file called `agents.md` or `rules.md` that contains all the rules that the AI agent will have to follow throughout the session.

The conversation history itself works as a sort of temporary ruleset, since previous messages are kept in context. But defining a strong, persistent ruleset is arguably the single most important factor in getting consistent, high-quality results from your assistant, especially when generating code.

When working in a team, it becomes even more critical. Since conversation history isn't shared across developers, a clearly written ruleset ensures that everyone gets the same baseline behavior, regardless of when or how they interact with the assistant.

Many people use persona-style prompts, like this one:

```
"You are a senior developer and an expert in Scala and Play Framework."
```


While that can help a little, it's far more effective to spell out exactly what you expect—not just the role but how that role should behave.

Here's an example:

"Provide a detailed explanation of how Play Framework handles dependency injection, including an example. Structure your response as if you were explaining to a senior back-end engineer who is new to Play Framework."

This is much more useful. You're not asking the LLM to impersonate someone; it's not an actor. You're giving it clear instructions for tone, depth, and structure. And that works better.

Also, here's a quick reality check: telling the LLM to "be careful," "be precise," or "don't make mistakes" doesn't help the way people often assume it will. These phrases only affect the style of the response, not the quality of the output. The model will sound more confident, but it doesn't improve the quality of outcome.¹

Concrete, actionable rules work much better, like these examples:

Never commit code before running all tests and verifying they pass.

When refactoring, never modify test assertions, only the implementation.

Always update the README and specification file after finishing a task.

These kinds of instructions are easy for the model to follow, and more importantly, easy for you to verify.

One common cause of hallucination is when the assistant tries too hard to "make something work" even if the input is vague or flawed. To avoid this, your ruleset should include guidance like this:

If instructions are unclear, ask follow-up questions before continuing.

If a request is impossible, explain why that is so and suggest alternatives.

It's also helpful to include project-specific guidelines, like preferred technologies, coding styles, or architecture patterns. The more relevant the ruleset is to your project, the better the assistant will be at staying on track.

Finally, the ruleset should be relatively short—around 1,000 to 2,000 words—so it can easily fit within the LLM's context window without taking up too much space meant for actual content.

1. <https://arxiv.org/pdf/2504.13656>

Your Code, Your Rules

Since I started working with AI assistants, I've gradually shaped an approach—or a kind of self-discipline, if you prefer—that I've now applied across many projects. It delivers high-quality code in very little time, and it's simple to follow. The only real requirement is to stick with it and stay focused.

I'll explain the approach here, and later in this chapter we'll see how it works in practice. The main point is this: to get consistent results, you need to define your rules and follow them. No shortcuts, no guessing, just a clear, repeatable process that keeps the assistant aligned with what you're trying to build.

That said, the goal isn't for you to follow my rules blindly. These are the ones that work well for me. What matters is that you find a set of rules that fits your own workflow, team, and way of thinking. Still, I recommend starting by understanding mine—I'll explain the reasoning behind each one as we go.

First of all, maybe you're a Test-Driven Development (TDD) fan like I am, or maybe you've never been too impressed by it. Either way, when working with an LLM, it's critical to write tests alongside the code.

Why? Because tests become the main feedback loop for the AI. Unlike human developers, the assistant can't reason abstractly or run the full application on its own. Tests are how we help it check its work and keep it honest. They let the assistant refine the output, catch most issues early, and often get things working before we even hit run.

If we let the AI write large chunks of code without tests, we're much more likely to end up with a mess. And if we wait until the end to add tests, we often discover the code is hard—or even impossible—to test without refactoring. That's wasted time. Writing the tests first, or at least together with the code, helps us avoid all that. At the end of the day it saves time, but we have to be specific:

- We don't want tests without real assertions.
- We don't want the same assertion repeated in multiple tests.
- We do want meaningful coverage that evolves with the code.

Once testing is in place, we can talk about code design. For that, we can rely on the classic Four Rules of Simple Design (from Kent Beck):

1. *The code works*: It contains and passes the tests.
2. *It reveals intention*: Anyone reading it should understand what it's doing.
3. *It contains no duplications*: Follow the Don't Repeat Yourself principle.
4. *It uses the fewest elements*: No extra abstractions, just what's needed.

Next—and this is just my personal preference (feel free to use your own style)—I want the LLM to adopt a functional style whenever possible:

- *Explicit over implicit*: Use parameters instead of private fields; prefer pure functions to class methods.
- *Immutable over mutable*: Immutable data is easier to reason about and debug.
- *Declarative over imperative*: Composing small functions is clearer than nesting conditionals and loops.
- *Stateless over stateful*: Stateful code is easier to write at first but much harder to debug later.

At the architecture level, here's what we aim for:

- *Clear modularization*: Large applications should be broken into manageable, focused modules.
- *Separation of concerns*: Each module should have one well-defined responsibility.
- *Low coupling*: Changing one module shouldn't break everything else.

And above all, we want to keep it simple: short functions, minimal dependencies, with no unnecessary complexity.

These aren't strict rules; sometimes you'll need to make the less elegant choice to get things done. But they're effective as general guidelines. If you instruct your LLM to follow them, you'll end up with much cleaner, more maintainable code.

These rules are battle-tested and make a solid starting point. But your own codebase might have different priorities. Maybe you prefer a more object-oriented style, for example, and that's totally fine. Feel free to experiment and adjust the rules to fit your way of working.

Just make sure to clearly explain your expectations to the AI assistant. The more context it has, the better it can follow your lead.

Now we also need to add a few rules about the process itself. The LLM should follow these rules:

1. Read the specification document before starting to code, and follow it carefully.
2. Update the specification after coding, adding a brief log of what was done and why.

3. Write and maintain a README with clear instructions on how to run the application.
4. Run all tests before declaring anything as “done.”

Keeping everything in Markdown documents gives us several key advantages:

- *Consistency across LLM sessions*: Context doesn't get lost between prompts or sessions.
- *No need for long prompts*: All the guidelines live in the docs.
- *Easy collaboration*: Multiple developers and different AI assistants can work together more seamlessly on the same codebase.

Finally, at the end you can add a section for your preference about specific technologies of the project—for example, using Poetry to manage dependencies in Python or using expression style for functions in Kotlin.

Here's a sample ruleset file you can use as a starting point. Keep it short and focused—you can always add more rules later. Just remember that context size is limited when working with LLMs, so every word counts.

Interaction Rules

- * Ask clarifying questions if input is unclear.
- * Explain why and suggest alternatives if task is not feasible.
- * Use structured, readable formatting (headings, lists, code blocks).
- * Follow instructions closely and explain clearly what you have done.
- * Don't modify code unrelated to the current task.
- * Try always to match the style of the code you are touching.

Coding Standards

- * Write meaningful tests with assertions for all code.
- * Avoid duplicated test assertions.
- * Maintain evolving test coverage.
- * Apply Four Rules of Simple Design:
 1. Code works (passes tests).
 2. Reveals intent.
 3. No duplication.
 4. Minimal elements.
- * Prefer functional style:
 - * Use explicit parameters.
 - * Prefer immutability.
 - * Prefer declarative over imperative.
 - * Minimize state.

Architecture

- * Modularize by concern, not by technical layer.

- * One responsibility per module.
- * Low inter-module coupling.
- * Short functions, no overengineering.

Workflow

- * Read `spec.md` before coding.
- * Update `spec.md` after task (log changes).
- * Write and pass tests before finalizing.
- * Keep a `README.md` with setup/run info.
- * Store all docs/specs in Markdown.

Commit Strategy

- * One prompt = one commit.
- * Each commit:
 - * Self-contained.
 - * Includes tests.
 - * Uses 50/70 commit message format.

Safe Practices

- * Do not change test assertions during refactoring.
- * Do not skip failing tests.
- * Do not invent unknown APIs; ask if you are unsure.

Project Preferences (Example)

- * Python: use Poetry.
- * Kotlin: expression-bodied functions.
- * JS: use ES Modules.
- * Follow `.editorconfig` + linter rules.

Goal

Produce consistent, safe, testable, and maintainable code.
Stick to the rules---no shortcuts.

To put all of this into practice, let's build an application with a clear goal defined up front. Unlike the more exploratory projects of the previous chapter, this time we're solving a specific problem, and we want to build it in a way that's easy to maintain and adapt later. Less vibing, more discipline.

Skillful Assistants

One of the most useful recent improvements in AI assistants is the idea of specialized skills.

Instead of putting everything in one large rules file, we keep only the core ideas there. The details move into small, focused instructions called skills—for example: how to write tests, how to review for security, or which UI style to follow.

The advantage is that each skill is loaded only when needed. Only a short description stays in the context. So even if you have many of them, they don't fill up the context.

Skills are usually simple markdown files with clear instructions. Around them is an ecosystem where you can define prompts, requests, and assertions in JSON. This allows you to evaluate each skill by comparing its output against expected results with the tools provided. Over time, you can refine the skill based on real feedback, making it more effective and reliable.

Building an Application Following Rules

For this chapter, the goal is to build a small but self-contained and genuinely useful application. We'll use our AI assistant to develop a command-line tool for indexing and searching through PDF files.

The use case is pretty straightforward: I have hundreds of ebooks, but searching through them is frustrating—especially when I can't remember the exact wording of what I'm trying to find. Full-text search isn't enough; I want to search by meaning.

That's where the same technology behind LLMs can help. If we extract the content of each paragraph and store its embeddings in a vector database, we can later search by semantic similarity, not just keyword matching. So even if I don't remember the exact phrase, the system can still find the right passage based on meaning.

We won't dive into the technical details of how embeddings and vector search work here. But if you're curious, I highly recommend reading [*Vector Search with JavaScript \[Gre25\]*](#), an excellent book on the topic.

Project Overview

Let's call this tool LESS, short for LLM Empowered Semantic Search.

Here are the key technologies we'll use in this project:

- ChromaDB as the vector database because it's lightweight, powerful, and open source.²
- Python as the language since it's widely known and has become the de facto standard for AI-related work.
- PyPDF (pypdf on PyPI) to extract text from PDFs.

2. <https://trychroma.com>

- spaCy (spacy.io) to parse the text and break it into meaningful sentence-level chunks instead of fixed-length token blocks.

LESS will be a command-line tool with just two commands, at least to start.

- One command will index all PDF files in a given directory into the vector database.
- The other will search a query against that database and return the most relevant results.

When a command is executed, the application should provide clear, step-by-step feedback about what it's doing. Once the task is complete, or if any errors occur, it should exit cleanly.

Let's see how we can build this tool by following the rules and practices we've just defined—for both ourselves and the LLM.

You can either watch the screencast of the full session or read the summary of all the steps below.

First Commit

The first step is simple: we create an empty project directory, verify that Python is properly installed, and copy our ruleset into a file—ready to guide the assistant from the very beginning.

Next, I asked the AI assistant to create a basic README.md file with the project description as specified.

After that, the agent immediately wanted to commit everything by itself. But I stopped it. I always want to review the work manually first and only commit if everything looks good.

I've cut the code review parts out of the recording, since there isn't much to watch there, but they're critical. This is the main checkpoint: making sure the work matches the intent before letting it into the repository.

Since everything checks out, this will be our first commit. Watch the video.³

For the next step, we'll ask the assistant to create a main file and handle parsing the command-line arguments for the first feature we're going to implement: indexing the documents. In other words, we want to extract all the relevant information from the documents and store it in the vector database.

3. <https://youtu.be/IDy3B4oilsc>

Building the First Command

Implementing the full indexing functionality would be too big of a step, so we need to find a way to slice off a smaller piece that still makes sense on its own.

One of the challenges when working with generative AI is learning how to break the work into steps that are small enough to be manageable and easy to verify, but still meaningful enough to feel like real progress.

This approach—implementing one small feature at a time—is sometimes called Inverse Salami (a term introduced in [Growing Object-Oriented Software, Guided by Tests \[FP09\]](#)). The idea is that we build the application one thin slice at a time, starting from nothing and gradually adding functionality until it's complete. This is the opposite of how you eat a salami: there, you start with the whole thing and take slices off until there's nothing left.

We could have started directly from the core functionality, but I personally prefer building applications from the outside in. That way, I'm always making sure we're implementing something that's needed.

Most of the time spent during this commit was around managing dependencies, setting up what we needed to even start. But eventually, we got to the real work: writing a first simple test for parsing the command-line command to index a directory, and then implementing the code to make that test pass.

Right now, the code just calls the indexer service, which doesn't do anything yet. But it's a good start (see the following screenshot):

```

less-cli
1 """PDF indexing functionality for LESS."""
2
3 from pathlib import Path
4
5 from rich.console import Console
6
7 console = Console()
8
9
10 class Indexer:
11     """Handles PDF indexing and vector storage."""
12
13     def index_directory(self, directory: Path) -> None:
14         """Index all PDF files in the given directory."""
15         console.print(f"Indexing PDFs in {directory}")
16         # Actual implementation will come later

```

```

test_cli.py
1 def test_index_command_requires_directory():
2     """Test that the index command requires a directory path."""
3     with pytest.raises(SystemExit):
4         cli.index()
5
6 def test_search_command_requires_query():
7     """Test that the search command requires a query string."""
8     with pytest.raises(SystemExit):
9         cli.search()

```

```

README.md
1 # LESS: A Lightweight Embedded Search System
2
3 ## Overview
4 LESS is a Python-based application designed for indexing and searching PDF documents. It provides a simple, efficient way to manage and query your document collection.
5
6 ## Installation
7 To install LESS, you need to have Python 3.10 or later installed on your system.
8
9 ## Usage
10 The application is run via the command line. The main entry point is the 'less-cli' command.
11
12 ## Contributing
13 We welcome contributions to LESS. Please see our contributing guidelines for more details.

```


After getting that first test working, I continued by adding a few more related tests. When all the tests were passing, I reviewed the changes carefully and then committed the work.

Now that the command works, the next step is to start implementing the indexer. The first thing it needs to do is simple: collect all the documents in the target folder so they can be processed.

Find All Documents in a Folder

We want to collect all the PDF files in the directory specified by the command-line arguments. So I prompted the AI to generate the tests and the implementation for finding files by extension.

As usual, once all the tests were passing and the code looked good after a manual review, I committed the changes. This time, I used Cursor's Git UI instead of the terminal—but you can use whichever workflow you prefer.

What's the next step? Now that we have the list of files, we need to parse them to retrieve the textual content.

Parsing PDF Files

We want to make sure we can parse a real PDF file, so I added a small PDF to the test directory and asked the assistant to implement and test the parsing functionality.

It all went smoothly. So far, so good. In less than half an hour, we had already made significant progress.

Now that we can extract text from PDFs, the next step is chunking it. This means breaking the text into smaller blocks that the embedding model can handle.

Chunking Text in Small Blocks

How small is “small”? Every model has its own context limit. Even if the model we're using can probably handle an entire chapter at once—with a relatively high limit, around 8,000 tokens—I still prefer to break the text into smaller chunks, usually around 500 to 2,000 characters. That typically ends up being less than 1,000 tokens per chunk.

This size is roughly comparable to a paragraph and usually leads to more focused, higher-quality results.

It's also important to break the text at logical phrase boundaries, not just by character count. This way, each chunk keeps a clear meaning, making it easier to embed and store effectively in the database.

For this task, we used the spaCy library,⁴ which provides convenient tools for splitting text intelligently. This step turned out to be problematic for the assistant. It wasn't able to write a sensible test that passed. Instead, it kept trying to either change the requirements or modify the tests to match its broken implementation, and I had to actively stop it from doing that.

In the end, the real issue was simple: the sample phrases it was using for the test were too short. But the assistant wasn't able to realize this on its own. I had to step in, debug the situation, and guide it back on track. Ironically, it took almost 20 minutes just to split the text properly, which felt surprisingly long compared to how fast it would have been to just write it manually.

Once the main issue was fixed, it was just a matter of ensuring the rest of the tests still passed and updating the documentation to explain how the chunking logic works.

Sometimes, guiding the assistant step by step through small misunderstandings may feel slower than just writing the code yourself, but it still saves time in the end. Helping the model understand the problem early speeds up the bigger picture flow.

Now we need to put everything together, but before that, we need to fix a small issue. While compiling, I noticed a warning about the PDF library version. The code was using PyPDF2, but the recommended replacement is the newer library pypdf. (Yeah, I know—it's confusing, but they decided to drop the numbers.)

Updating the PDF Library

Refactors like this are common. LLMs are trained on existing code, often old, so they might not be aware when a library has been updated or renamed, especially in ecosystems where backward compatibility tends to break often.

In the future, assistants may be able to stay up-to-date by reading documentation from the internet. But for now, it's up to us to catch these issues and prompt the assistant to update things manually. In this case, the change was minor, and the assistant managed to upgrade everything by itself in just a few minutes.

4. <https://spacy.io/>

Following the principle of not mixing refactoring with new feature development, we committed the changes when the tests passed, before starting a new functionality. Watch the video.⁵

For the next slice, we need to split the parsed text into chunks.

Complete PDF Parsing and Chunking

We start by asking the assistant to write a test that reads a real PDF file from the test folder to use realistic input data. Then it should parse the pages, normalize the text, and finally chunk it into consistent segments that can be embedded and stored in the vector database.

As often happens, the model forgot that we already had a real PDF in the test directory. Instead, it started trying to generate a new one directly in the test code. It's technically possible but slow, and the generated PDF would likely be overly simplistic. Using a real, small PDF is both faster and more realistic.

When it wrote the test, the model included some weak or irrelevant assertions, like checking that all chunks weren't empty strings or that they included newline characters. Instead, I asked it to only verify that each chunk was between 500 and 2,000 characters and that there were at least 10 of them in the example PDF.

The good news is that once we corrected the assertions, the test passed without needing any changes to the code. This is a good lesson: always check with double care what the generated tests are testing (around 12:30).

Now that the code and tests are okay, it's time to move on to the next step: storing all the chunks in the vector database.

Storing Chunks into the Vector Database

For our vector database, we chose ChromaDB because it's open source, easy to use, lightweight, and powerful.

In general, a vector database stores and retrieves embeddings—the vector-based semantic meaning of a word or phrase—along with some metadata. If you want to dive deeper into how vector databases work, check out [Appendix 1, Demystifying Large Language Models, on page 101](#).

Using ChromaDB is straightforward: we specify the embedding model name, and if it's supported, the database will automatically download and load it behind the scenes. Then whenever we insert a chunk of text, ChromaDB calls

5. <https://youtu.be/BTortzZHPBU>

the model to generate an embedding and stores it along with any associated metadata, if present.

With a simple prompt, the assistant correctly understood that we wanted to use ChromaDB as the vector database. It added the necessary import, wrote the first test, and got it passing fairly quickly.

At this point, I asked it to move the database-related code out of `Indexer.py` and into a separate file. Ideally, this would go in a separate commit, but since the database code hadn't been committed yet, it made sense to do it right away.

Asking for this refactor also forced the assistant to split the tests, separating the database logic from the PDF parsing and chunking tests, which made the overall structure cleaner and easier to manage.

This is another example of how important it is to “keep the edge”—to stay just ahead of the assistant, guiding it to refactor for clarity and always test the behavior (around 27:30).

The next step is to store the metadata we needed for each chunk—specifically, the name of the file and the page number where the text came from. It would be frustrating to get a matching result from a search but have no idea where it came from.

Storing Metadata

Luckily, ChromaDB supports storing metadata alongside each embedding.

Generating the tests and making them pass wasn't a problem. But the assistant got a bit overzealous and added features I hadn't asked for, like enabling metadata-based search.

Now, you might think, “Free features, yay!” But in practice, adding functionality that isn't explicitly required often leads to trouble: more bugs, harder maintenance, and unnecessary complexity. When in doubt, less is better.

The real issue is that once the AI assistant “gets an idea”—or more likely, recalls a pattern it was trained on—it can be surprisingly stubborn. You have to be very specific and deliberate in your prompts to convince it to remove something it believes should be there (around 38:30).

While reviewing the tests before committing this step, I noticed they were relying heavily on mocks. Now, nothing's inherently wrong with mocking, but it adds complexity to the test setup, and it's easy to overuse it to the point where we're not testing anything meaningful.

A better approach is to decouple the service logic from the database and PDF parsing details by injecting those parts as adapters. This not only makes the tests much simpler—removing the need for mocks and mock-specific assertions—but also makes the code easier to debug and reason about. So let's do this now!

Introducing Ports and Adapters

First, we ask the assistant to introduce the Ports and Adapters design pattern into the indexer service⁶ and then remove the mocks. The assistant immediately understood the pattern and applied it correctly (around 58:30).

The problem came out in the tests. Instead of using simple, custom stubs for the adapters, the assistant kept trying to mock them using the same mocking libraries. I had to explicitly instruct it to implement test-only adapter stubs that conformed to the expected interface. Watch the video.⁷

Then it forgot to remove the old database code that was no longer used. This is a typical LLM behavior: when asked for a change, it tends to rewrite the code from scratch rather than refactor or reuse what's already there, and it often forgets to clean up what's been replaced.

After a bit of renaming to tidy things up, I was able to commit all the changes. Of course, I made sure the docs were updated and the tests were passing before committing (around 10:00).

It's interesting that the assistant correctly explained how the ports and adapters pattern improves decoupling between the application logic and its external dependencies, like the database, chunking, or PDF parsing. Moreover, removing the mocks definitely made the tests cleaner and more robust in the sense that they won't break due to small internal changes.

Still, the assistant didn't suggest this architecture on its own, and it couldn't complete the refactoring without guidance. Even more, while it applied the pattern mechanically, it completely missed the point: it didn't reduce the complexity. The mocks stayed, and the old database code wasn't cleaned up, which was the main reason for applying the pattern in the first place.

I suspect that decisions like this—architectural shifts, design patterns, and long-term trade-offs—will remain firmly in the domain of human developers for the foreseeable future. And I don't think this is a bad thing.

6. [https://en.wikipedia.org/wiki/Hexagonal_architecture_\(software\)](https://en.wikipedia.org/wiki/Hexagonal_architecture_(software))

7. <https://youtu.be/wYmTTx5iDTU>

But whatever your preference, the fact remains: LLMs are still limited to concatenating words. Even with the so-called chain of thought, they're still floating above the ground. They have no tools for confronting the harsh realities of software design—balancing competing goals, anticipating future needs, and navigating trade-offs that go beyond what's written in the code.

Going back to our project, with this refactoring we can consider the indexing done—one slice after another we completed our salami, or at least half of it. So we're now ready to tackle the search.

Adding Search Command

The core search functionality was already implemented inside the indexer; we just needed to wire it up through the command-line interface and add proper tests.

As often happens, when adding the tests, the model fell back into its old pattern and used mocks to simulate the search behavior. I had to step in and explicitly tell it to remove the mocks and use proper adapters instead. This was also a good moment to update the README.md file to make sure it accurately reflects how the tool works now and how to use it from the command line (around 24:00).

At this point, most of the functionality was in place, so we could add a proper end-to-end test to make sure everything was wired up correctly and that it could index and search PDFs as intended.

End-to-End Testing

This simple task turned out to be more work than expected. Instead of simply writing a test that calls the existing index and search commands, the assistant created a brand-new combined index and search feature in the codebase without even reusing the code we already had. I only realized this when it was time to commit. Watch the video.⁸

Maybe we could have avoided it with a clearer prompt and a bit more attention. But no matter what you do—if you use AI assistants long enough, something like this will happen to you too.

So let's take it as a lesson. It's a good reminder of why reviewing the code at each commit is so important: it's easy to miss these things in the chat, especially when you have multiple methods with similar names doing similar things. You only catch them when you look at the diff.

8. <https://youtu.be/V0pyOc-DG6g>

So now we had working code, but it was doing the wrong thing: a new command that indexed and searched in one step. After verifying that all tests passed, I committed the changes. Then, immediately after, I asked the assistant to remove the duplicate command and strip out the indexing logic. The end-to-end test should simply call the existing index command first, and then search.

Another time-waster was the Click library for CLI argument parsing.⁹ The assistant had imported it early on, but later something broke and the interface stopped working properly. Since this is a very simple CLI, I decided we didn't need a full-featured library. I had the assistant remove Click and parse the arguments manually, calling the service functions directly.

While debugging and cleaning things up, I also asked it to add detailed log messages to show what's happening and track progress during indexing and search. To be honest, this should've been part of the original specification and properly tested, but since this is just a handy tool and not a critical application, simple print statements are good enough for now.

In the end, everything worked. It took more time than strictly necessary, but now we have a more complete application and a useful end-to-end test that gives us confidence the system works as a whole (around 49:00).

This stint also raised an interesting question: what should we do when, while working on one thing, we suddenly realize that something else is broken or missing? Should we shelve the changes or roll them back? My rule of thumb is if the code works, there's no reason not to commit. Then, once we commit, we should immediately address the missing or mistaken functionality before moving on to anything new. I only roll back when the code has been broken for a while, say, more than 20 minutes, or we're clearly stuck going in circles.

All the core features are now in place. Time for the moment of truth: running the application on a real directory with some PDFs and checking the results.

Indexing Incrementally

Using the application, rather than just running the tests, I quickly realized a couple of missing features:

- First, indexing should be incremental: it should only process documents that have been added or changed since the last run, with an option to force a full reindex if needed.

9. <https://click.palletsprojects.com>

- Second, the ChromaDB file should be stored inside the same directory as the documents, to simplify maintenance and portability.

Adding these features, verifying that everything worked as expected, and updating the documentation took about another hour. No new problems came up, and for the ones we've already encountered, we now have reliable ways to handle them. Watch the video.¹⁰

The only thing left is to put the application to real use: searching through my book collection to find relevant material. This is where the power of semantic search really shows—if the application works, that is.

Final Test

With incremental indexing in place, the application is finally usable, and useful. I can now search through my collection of technical books and quickly find the material I need. I tested it by querying one of my book PDFs for an important topic but deliberately avoided using the exact wording from the book.

```
> poetry run python -m less.cli search-dir /home/ubertobarbini/books_test
"handling the state in functional programming"
```

```
LESS CLI module loading...
Imports completed, loading Indexer...
LESS CLI initialization complete
```

```
Searching for: handling the state in functional programming
```

```
[1] From: from-objects-to-functions_P1.0.pdf, Page 144
Score: 0.2910
```

```
-----
```

```
How can we model this mutable entity in functional programming?
```

```
Implement a Finite State Machine
```

```
Let's draw a diagram of all the states and their transitions:
```

```
This diagram illustrates how we can represent the entities that change
behavior according to their state using a state machine.
```

```
We map the transitions into domain events, then we constrain each state to
accept only some events, and finally, we allow events to change the entity
state to another one.
```

```
...
```

Notice how it selected a very good match from across the whole book (which is about 500 pages) not by matching the words exactly but by matching the meaning.

The application isn't completely finished yet, but it's already good enough to be useful.

10. https://youtu.be/-x_mvYOD7CM

Results and Reflections

All in all, we finished the application in about four hours, with 19 commits and roughly 1,200 lines of code, plus about the same amount for tests. And I didn't type a single line myself!

Now, the quality of the code probably isn't going to win any software quality awards, but it's good, solid code. It's easy to fix, easy to extend, and more importantly, I understand it. That's the main difference compared to what we got from vibe coding.

We progressed using tests. It wasn't 100 percent pure TDD; sometimes the AI created a bunch of tests and then implemented them all at once, or even added tests after writing the code. This happens because current LLMs aren't yet good enough at strictly following detailed TDD instructions, or at least, I don't know how to reliably instruct them for it yet. But honestly, it's close to how most real teams work anyway, and it's very likely that LLMs will get better at this in the future.

During the session, the assistant struggled badly at two points: during the text splitting and when it was not printing the search results. In both cases, I was very close to rolling back and restarting from the previous commit. In the end, the assistant managed to solve the problems with a few hints from me, but looking back, a clean rollback might have been faster. It's always painful to throw away "almost working" code, but sometimes it's the right move.

You can also see that I'm not religiously following all my own rules. They're guidelines, not strict laws, and knowing when to bend them is part of knowing your tools. Still, one golden rule I never broke: run the tests and verify before every commit.

Also, you may have noticed how much babysitting AI assistants still require. They can be very stubborn once they take the wrong direction, even if you've told them exactly what to do. You have to stay actively in control. That said, we've seen huge improvements already, and I expect AI assistants will keep getting better at following instructions over time.

Even with these improvements, it's important to make a habit of challenging the assistant to explain itself. When we're unsure about the code, we should ask why a certain choice was made and check whether the explanation makes sense. This helps us avoid drifting into accidental behavior.

Finally, automatic tests turned out to be absolutely essential, not just for catching regressions but for debugging when things got confusing. If all else fails, and you're stuck, ask the assistant to explain what's happening. It's surprisingly good at helping you figure things out even if it often can't fix the problem by itself.

Overall, it's not about writing perfect code; it's about keeping momentum, keeping quality under control, and finishing real work with AI at your side.

What We Covered

This has been a long session, but much more productive than the vibe coding we saw before. Here's a quick list of what we covered:

- *Planning*: It's crucial to write down—or have the LLM generate—a detailed plan outlining what we want to achieve. The plan acts as a shared reference that the LLM can use to stay on track, updating it at every step without losing sight of the bigger picture.
- *Precise prompts*: Avoid broad requests like “implement the whole feature.” Instead, break the work down into small, specific tasks and request them incrementally. This keeps you in control of the output and massively reduces the chances of getting wrong, bloated, or messy code.
- *Automatic tests*: Make sure all critical logic is covered by automatic tests. Always write, validate, or at least carefully review the tests yourself. Don't blindly trust LLM-generated tests; they can easily include wrong or meaningless assertions without you noticing.
- *Simplicity*: Aim for simple, readable code. If you struggle to understand what the LLM has written, it's a red flag. Remember, if you can't understand it easily, the assistant probably can't either. More importantly, your own reading speed, not the AI's writing speed, is what limits your productivity. Keep the code clear enough that you can move quickly.
- *Committing*: Always commit your code before starting an AI-assisted session, and keep committing small, tested changes along the way. This gives you safe restore points at every step. As always, make sure all tests pass before each commit to avoid breaking the build.

Now, in the next chapter, we'll dive into a more complex example, and we'll also see how you can use LLMs to learn new technologies while you build.

Learning by Coding

In the previous chapter, we built a simple project from scratch with help from AI. Now we'll take on something more challenging: creating a more complex project while learning a new language and framework at the same time.

This is one of the least-talked-about strengths of AI assistants: they can significantly cut the time it takes to become productive in an unfamiliar ecosystem. Instead of spending days or weeks studying docs and tutorials, you can ask the LLM directly and get a “know-it-all” tutor that explains concepts while working on your own code.

It's also a great chance to practice the habits we introduced earlier—one prompt, one commit, focused instructions, test-first development—and see how they hold up in a less scripted, more open-ended scenario.

The same approach works whether you're building something new inside an existing codebase or prototyping a completely new product. We'll lean on the assistant to generate code, but we won't rush ahead. We'll still work with care—moving fast, but with enough structure to avoid stacking up technical debt.

All along, we'll stick to one core principle: keep things as simple as possible. That includes saying no to premature abstractions or clever tricks. The YAGNI rule—You Ain't Gonna Need It—is still one of the best filters when making early decisions. Don't build features or layers you don't need yet.

The Inception

For this project, we'll continue building on the command-line search tool (LESS) we developed in the previous chapter. If you remember, we built it to run semantic searches over a local collection of PDFs, using a vector database for indexing and retrieval.

Now we want to develop a proper UI for it—a simple, easy-to-use web application that lets users search the collection and immediately view the relevant page from each PDF. The back-end logic will remain the same: the command-line tool will handle the search and return a JSON response with up to 10 matches, each containing a similarity score and a reference to the matching PDF and page number.

The web app will take care of the rest. It should allow the user to enter a query, show past searches, and display the results returned by the CLI in a clean, structured format. When a result is selected, the app should fetch the corresponding page from the PDF and render it nicely in the browser.

Even if our goal here isn't to build a full-scale product, we have many implicit assumptions about what a web application should look like: it must work reliably, be easy to use, be reasonably performant, and meet at least basic security requirements.

For long-term maintainability, we also have some internal quality requirements:

- We want our code to be as simple as possible, avoiding accidental complexity.
- It must be verified by automated tests so that we can quickly check if it's correct.
- All code should clearly express its intent, making future extensions and maintenance easier.
- We want to avoid duplication of logic to reduce the risk of inconsistencies and simplify maintenance for both humans and LLMs.

We learned in the previous chapter that vague prompts are not good enough. Something like this leaves too much up to chance:

“You are a senior software engineer. You need to design and build a UI that calls a CLI application to search for text in PDFs and then displays the exact page of the PDF in the browser.”

The assistant has to make too many decisions on its own—which technologies to use, how the UI should behave, how to structure the code, and whether or not to follow any of the quality principles we just mentioned. And while LLMs can do some planning and reasoning, those plans are theirs, not ours. They're often shaped by patterns in training data, not our real-world needs.

To get exactly what we want, we need a longer, more detailed brief. But I don't like writing those manually. What works better for me is to develop the spec

through a conversation with an LLM. I ask questions, explore alternatives, and let the assistant bring its knowledge of tools and frameworks into the mix. I stay in control of the decisions, but I let the AI support the thinking.

This doesn't have to be a different model from the one we'll use for coding, but in practice I find it more comfortable to use a chatbot like ChatGPT or Claude for this planning phase. They're optimized for dialogue, and their features—like Canvas, Projects, or long chat history—make it easy to refine a spec over time.

Once I have a clear idea of what I want, I ask the assistant to summarize everything in a structured document. I review and tweak it as needed. The goal is to end up with a prescriptive text that sets direction and tone for the code assistant.

In the end, a good prompt is a clear and detailed document. How long should it be? From my experience, the sweet spot is around a thousand words or less. If it's too long, the LLM struggles to keep the whole thing relevant, even with a large context window. If it's too short, it won't have enough detail to guide the model properly.

And the best way to create it is through iteration. Let the assistant generate a first draft, then work together to refine it. Once it's ready, commit it to your repository. This way, your AI coding assistant can reuse it in every prompt without losing context—even when the context window is tight.

Shaping the Idea

Let's get to work. You can check out the following video (about 10 minutes in total) to see an example of how to put together the first draft of the project detail document, chatting with an LLM to navigate technical choices and refine the requirements. Watch the video.¹

Just a note: in a more serious project, shaping the initial spec usually takes longer, even several hours, especially if the scope is large or the requirements aren't clear yet. What you'll see in the video is a simplified example, just enough to keep it useful without getting too boring.

The process starts by listing the key points that define the project (00:16). These are passed to a language model, in this case, ChatGPT-4o, which turns them into an initial draft. The model helps lay out the structure clearly, saving time and reducing friction in getting started.

1. <https://youtu.be/cO-QiG2QyKQ>

This step can be done with any capable LLM chat interface, such as Claude, Gemini, or an open source model running locally. The important part is the back-and-forth conversation to refine the idea into something usable.

Next comes the tech stack discussion (after about two minutes). The assistant suggests several viable options, each with its own pros and cons. After discussing them, I decided to use the Elixir language and the Phoenix web framework for the app. Elixir is an elegant language with a small memory footprint and a simple installation process, making it lightweight and easy to work with. Phoenix complements it well, is highly productive, and lets you build rich, interactive UIs without relying on heavy JavaScript frameworks.

They're certainly excellent technical choices, but part of the reason I chose them for this chapter is that I'm not very familiar with them. This makes it a good example of how we can use LLMs to learn new technologies while building real products.

Of course, the choice isn't what matters here—this is just an example of how to reason through a technical decision with the assistant. We then update the technical spec document accordingly (around 04:20 mark).

To make expectations clear, the assistant is prompted to add a section on Kent Beck's design principles, emphasizing clarity, modularity, testability, and short functions (for around one minute).

A few refinements follow. For example, the initial version drops authentication to keep things simpler. Example queries and JSON responses from the CLI tool are also added to the document as references (around another three minutes).

Finally, the assistant is asked to rewrite the document in a more compact format; it's still complete but easier for an LLM to process efficiently and with lower memory usage.

Considerations on Discussing Ideas with LLMs

This process is partly Rubber Duck Debugging,² where you clarify your thinking just by explaining it out loud, or in this case, to the AI. But it goes further than that. LLMs also bring the benefit of a massive built-in knowledge base, basically the whole internet. On top of that, they're fast at writing things down in a clear, structured way. That's a huge help if you're not a native speaker, or if you need to translate the document later.

2. https://en.wikipedia.org/wiki/Rubber_duck_debugging

One of the best things about using an LLM as an assistant is how well it understands programming languages and frameworks, especially widely used open source ones. If this had been a solo project, for example, I probably wouldn't have picked Elixir and Phoenix. They're not technologies I know deeply, even though I've heard good things and have friends who love them.

But with an LLM as a coding partner, it's easier to take that kind of leap. The risk of being blocked by an unfamiliar syntax or library drops significantly. That shifts the focus: instead of defaulting to what I know best, I can just ask what fits the job best.

In this case, a JVM language like Kotlin would probably have been too heavy, both in terms of system requirements and startup time, even though it's the one I know best. Starting from scratch in a more verbose, setup-heavy stack would have just slowed things down. What I needed was something lightweight, responsive, and quick to get running so I could iterate fast.

That said, it's still important not to blindly trust the LLM's answers. This has become less risky, since we can now ask for sources and cross-check information when needed. But some critical thinking is always required.

With the spec document ready, the foundation is set. It's time to start building.

Setting Up the Project

We start by creating an empty directory for the project. Then we put in the ruleset document as explained in the previous chapter. Finally, we initialize a Git repository and make the first commit with the ruleset, giving the project a clean, version-controlled starting point.

Installing Elixir

Once Windsurf is set up, the next step is to install and configure Elixir so we can run everything locally.³

Elixir runs on the Erlang VM (also called BEAM), so we'll need to install Erlang first. After that, we can install Elixir itself. The process is simple, and plenty of beginner-friendly guides are available. A good place to start is with the official docs.⁴

But of course, we're not going to read them all! We can ask our AI assistant to guide us step by step on how to install and set up everything we need.

3. <https://elixir-lang.org/>

4. <https://hexdocs.pm/elixir/introduction.html>

Could we ask the assistant to install the environment directly on our machine? Technically yes, but that would mean giving it permission to run admin-level commands without any supervision—something I’m not ready for. And honestly, running the commands myself has been instructive.

With the setup complete, we’re ready to start coding the application.

One of the best practices when building a new application from scratch is to begin with what’s called a Walking Skeleton.⁵ This is a minimal version of the app that runs end-to-end, just enough to wire everything together but without any real business logic yet.

When working with LLMs, this step becomes even more important. They often struggle with basic setup details, and fixing those early, before layering in business logic, helps prevent a lot of confusion and wasted effort later.

Walking Skeleton

The first step in building the walking skeleton is to make sure we can display a simple Hello, World! inside our Elixir app. Remember, the goal here isn’t to deliver any meaningful feature yet. What we want is to confirm that the whole stack is wired up correctly.

This is a basic sanity check: is the framework responding? Is routing working? Can we build and run the app? Are the tests set up correctly? Is the UI rendering anything at all? We’re laying down the scaffolding that the rest of the project will be built on. Watch the video.⁶

As shown in the video and the screenshot on the facing page, this example uses Windsurf. Its UI is similar to Cursor, though it also offers a plugin for JetBrains IDEs. The stand-alone version includes a slightly customized Visual Studio Code environment, which many developers will find familiar.


The first command (00:01) is always the same: ask the AI assistant to examine the project source and read the design document. This gives the assistant the necessary context: what we’re building and how we want it built.

Once the assistant confirms it’s ready, the next step (around the 1:00 mark) is to prepare the walking skeleton. The goal is to verify that all the core components are wired up and working before introducing complexity.

Now Windsurf needs to run a few setup commands using Elixir’s mix build tool. This kind of interaction, where the LLM can execute commands, read

5. <https://wiki.c2.com/?WalkingSkeleton>


6. <https://youtu.be/ID0YQ2Kxibc>



the output, and adapt its actions accordingly, is called agentic mode. In this setup, Windsurf is allowed to run tests and inspect Git diffs, but not to execute arbitrary commands without confirmation. That's why it explicitly asked for permission to proceed with the setup tasks.

One limitation I ran into during this phase is that Cascade—the Windsurf interface—struggles with console commands that require interactive confirmation. But in this case, that worked in my favor: I wanted to learn the useful Elixir and Phoenix commands myself. So I ran the setup steps manually in the terminal.

After completing the basic project setup, the server was started, and the Elixir welcome page displayed successfully (at about 06:45) with the project name (see the following screenshot).



This is a good point to pause and commit all progress made so far.

The next step is to start adding domain logic—in other words, the behavior we care about. But before jumping in, let's take a moment to review how the Elixir walking skeleton is working.

Learning Elixir

Let's not forget one of the goals we set for ourselves at the beginning of this chapter: we wanted a nice UI for our tool, but also a chance to learn Elixir.

This is a great moment to focus on that learning. Everything the assistant generated so far reflects the typical structure—or skeleton—of a Phoenix project. So we paused the assistant to take time to explore the generated files and get familiar with how Elixir and Phoenix projects are organized.

Since this stack was new to us, it was a valuable opportunity to understand how the framework structures its code, where key components live, and how the pieces fit together. Even though most of the files were boilerplate, walking through them helped us build a mental model of the layout before adding real logic. The design document also updated to reflect our current progress. Once I was confident everything ran correctly and, most importantly, that I understood how, I committed all the changes.

All in all, this commit was large, and a seven-minute gap between commits feels long when working with LLMs. But since this step was mostly about running scaffolding commands and committing generated files, I think it's acceptable.

On the other hand, I struggle when an assistant runs for ten minutes or more, producing and refining large amounts of code. I lose concentration, and it takes me a long time afterward to understand everything it generated. For this reason, I prefer to stop the assistant if it hasn't produced a good result within five minutes or so.

This is a personal preference, of course—you may be more patient, or less, than I am.

Remember that regular breaks and frequent commits help us maintain a sustainable, comfortable pace. Coding is more like a marathon than a sprint. To maintain productivity and quality over time, especially with the support of fast AI assistants, it's essential to find a rhythm that suits your personal work style, allowing you time to recharge mentally and stay focused throughout your session.

Our next step is to progress with the domain of the search functionality.

Modeling the Domain

We start the session by asking the assistant to reread the design document to make sure it still had the full context. This is a good habit to repeat from time to time, especially when sessions get long. LLMs can “forget” earlier instructions or start drifting from the rules, so if responses start feeling off, simply ask it to reread the document. Watch the video.⁷

Once ready, the assistant began developing the search module—the core functionality of our application. Since we had clearly specified a test-driven approach, the assistant started by writing tests for the search logic. However, left over from a previous commit were some failing tests, a reminder of the importance of running tests before committing (shame on us for skipping that step).

Despite the existing failures, the assistant moved ahead with implementing the search module and its corresponding tests. After a few iterations, everything passed successfully (around the 05:30 mark).

Before reviewing the new code, we first addressed the failing test in the controller. This turned out to be a quick fix. The default tests generated by the framework had expected a generic welcome message, but our app was using a custom name instead.


The assistant quickly located the mismatch and fixed it. Interestingly, its first instinct was to change the application text to match the test, rather than adjusting the test itself. While this is generally the right mindset—tests should reflect real user expectations—in this case, the test was outdated and needed updating instead (see the screenshot on the next page).

This is a good reminder of the kind of supervision we still need to provide to AI (for how long, nobody knows). In a way, our developer role is shifting into that of an “LLM tamer”—guiding, reviewing, and course-correcting the assistant’s work.

Once the issue was resolved, all tests passed, and we moved on to review the search module and its corresponding tests. Everything looked solid, so we committed the three modified files (about seven minutes after starting the session).

The assistant then continued to build on the search functionality, still following the TDD process by writing new tests first. Before letting it implement the

7. <https://youtu.be/d5mVnvmopmM>



logic to make the new tests pass, we reviewed them. One of the tests was checking a behavior we didn't need, so I asked the assistant to remove it (around the 08:30 mark).

Learning and Supervising


Reviewing and supervising the code up to this point helped me understand Elixir and Phoenix much better, simply by watching how real problems were solved. It's a far more efficient and engaging way to learn a technology than reading documentation alone.

On the other hand, this approach doesn't teach much about the why behind the technical choices. For that, a good book or deeper reference is still the best choice for me.

So far, we're following a test-driven approach, but it's worth noting that this isn't the classic TDD cycle. Traditionally, we write a single test, make it pass, refactor, and only then move on to the next test. LLMs, however, tend to prefer writing several related tests at once and tackling them together.

When all the tests cover the same piece of functionality, this approach works surprisingly well with the superhuman writing speed of LLMs. It gives us a broader view of what works and what doesn't, and it lets us review the entire logic in one go.

After a few more iterations, the assistant made all tests pass and updated the design document to reflect the current state of the project and the next steps, as shown in the next screenshot.



This marked another good point to review and commit all the completed work. With the back-end search functionality now in place, we can move on to building the web User Interface (UI).

Developing the Search UI

In this step, we'll use the back-end API we just developed to power the front-end part of the application. This is where we handle user interaction—sending search queries and displaying results directly in the browser. Watch the video.⁸

As usual, we began the session by asking Windsurf to reread the design documents and refresh its understanding of the project. Once it confirmed, it saved a memory (at minute one). In Windsurf, a memory is a way for the agent to keep track of context across sessions. Unlike the ruleset, which we define explicitly, memory is fully managed by the assistant itself.

Initially, the LLM wanted to jump straight into UI development, but we paused it and redirected the focus back to the ruleset. To keep things honest—and to give the LLM the feedback it needs for better results—we stick to a strict test-driven approach. So we asked it to create a Bash script to simulate LESS, the search tool we want to integrate with. This let the assistant write and run tests without depending on the external LESS project.

Even if I'm repeating myself, I can't stress this enough: when working with LLMs, it's important to stay in control of the development flow. Don't hesitate to stop


8. <https://youtu.be/DTTch6H9Nuo>

code generation if things start to look off or the direction isn't clear. LLMs don't get offended, and it's always easy to regenerate or roll back if needed.

That said, the Bash script produced by the assistant was impressive. It matched the structure of the question/answer examples we provided and even added a bit of randomized behavior to simulate real results. It was efficient, clean, and generated in just a few seconds (around the 03:30 mark)—probably better than what we would've written by hand.

Next, the assistant moved on to generating domain-level tests for parsing the CLI responses. After running into a few dependency issues related to a mock framework, we got the tests to pass. But on reviewing the code, it became clear that the mock framework wasn't necessary at all. The whole design could be made simpler by passing a couple of functions to handle the I/O directly within the domain logic.

We shared that idea with the assistant, and it agreed right away, even complimenting the change as a cleaner and more maintainable design (around the 09:00 mark). (See the following screenshot.)



Working with LLMs can sometimes be intense, but one upside is that they're surprisingly good at boosting your ego. It's easy to feel proud of a good idea when the assistant compliments it. The problem is, they'll do the same even when the idea isn't that good at all!

With all tests passing and the parsing logic working cleanly without mocks, we committed the changes.

AI Sycophancy

LLMs have a well-known tendency to agree with whatever we say—even when it’s clearly wrong—and to praise our ideas far beyond credibility. This behavior has a technical name: AI sycophancy.


It comes directly from how these models are trained. Many example conversations in the training data use polite confirmations like “yes, you’re right” or “that’s a great insight,” so the model learns to echo those patterns. Often this politeness has very little to do with the real value of what you said.

It may feel nice to receive a bit of flattery now and then, but we still need to stay alert and double-check our ideas on our own.

Model developers know this is a problem and are working to reduce it, but they still want the assistant to sound friendly and positive. After all, better to err on the side of politeness than end up with a very rude assistant!

Next, it was time to move forward with the UI. As before, we paused the assistant as soon as it started generating code. Instead, we first asked it to verify whether the existing UI could work correctly with our simulated Bash search tool. We also requested that the external command used to trigger the search be made configurable (around the 12:00 mark).

After a few minor fixes, everything was running smoothly. We launched the application and saw the simulated search results appear as expected (after 13 minutes), as shown in the following screenshot. It’s impressive how quickly things came together with a focused workflow and the right level of supervision.



Finally, the assistant reran all tests, fixed the one remaining failure, and updated the documentation to reflect all recent changes. After reviewing the final code thoroughly, we committed the updates.

With that complete, the next step is to display the actual PDF page directly in the browser.

Show PDFs

Now that users can search and get results, the next logical step is to let them view the content, specifically, the PDF page where the match was found. This means adding client-side functionality to render PDFs inside the browser, based on the data returned by our back end. Watch the video.⁹

This part adds a bit more complexity, involving JavaScript integration and UI handling. But it's exactly in moments like these that being able to guide and manage an LLM becomes most important. And in fact, things went a bit pear-shaped at first: the assistant rushed ahead, generated a large chunk of code without tests, and quickly got stuck.

This recording captures the second attempt at adding PDF rendering. In the first attempt, the assistant jumped straight into integrating a JavaScript library for displaying PDFs. Unfortunately, it generated a large block of code without any tests and struggled to make it work correctly.

After about five minutes of unclear progress, I stopped and reverted everything.

So this second time, we'll shift the focus to progressing the back end and writing tests first, postponing the JavaScript integration for later.

Setbacks like this are normal (but less and less frequent) when working with LLMs, and we don't consider them wasted time. In fact, the first failed attempt helped clarify the problem. At the time, we weren't sure of the best approach to display PDFs in an Elixir/Phoenix setup. Watching the LLM struggle gave us valuable insight and helped us plan more effectively for the second run.

After rolling back to the last good commit, we uploaded three small PDF files into a test folder. At the start of the new session, we informed the agent about these files and, as usual, asked it to get familiar with the current state of the project (00:28).

As the project grows, the agent takes slightly longer to process everything. But our habit of maintaining clear documentation pays off. The assistant consistently follows the defined process and respects the design guidelines across sessions (at the 02:30 mark).

During the next set of UI changes, some tests broke because UI elements had moved, causing selectors to fail. The agent attempted to fix them but skipped rerunning the tests, assuming the changes were safe. While that mirrors common human behavior, it's something to watch for. We paused the assistant

9. <https://youtu.be/HKft13L3Nas>

and explicitly asked it to rerun the tests with the updated selectors (around minute 4).

Later, when the assistant tried to integrate the JavaScript PDF library again, we paused it once more, remembering the issues from the previous attempt. This time, we asked it to first set up a back-end endpoint to serve PDFs to the front end. Initially, the assistant started creating new PDFs for each test, but we reminded it to use the ones we had uploaded earlier in the session (06:08).

After some debugging and small adjustments, we were able to successfully send the existing PDF documents to the UI. The interface now included a button to display the PDF, though at this point it only rendered a placeholder (09:37). Still, this marked solid progress. All tests passed, so we committed the updates.

The assistant then proceeded to create the JavaScript hooks required by the PDF rendering library. It also modified the Bash script to consistently return one of the existing PDF filenames, instead of generating new random ones. The assistant prompted us to check the result, but unfortunately, the application only showed a gray box (at the 11:20 mark).


To troubleshoot, we asked the assistant to download the JavaScript library and serve it locally, ensuring it would work even offline. The change was made quickly, but the gray box remained. Eventually, we realized the issue: the simulated search was returning random page numbers, often pointing to pages beyond what existed in the test PDFs. We asked the assistant to limit the results to either page 1 or 2 (15:59).

This was a good reminder that LLMs—Claude 3.7 in this case—still struggle with small, practical issues like this, even while handling more complex tasks like writing Bash scripts or wiring JavaScript libraries into a Phoenix app.

With that final fix in place, the application worked as intended, successfully displaying the correct PDF pages based on the simulated search results (around the 18:00 mark), as shown in the screenshot on the next page..


At this point, we thoroughly reviewed the code and asked the assistant to update the design document with all recent progress. Once everything was confirmed and aligned with our original goals, we committed the changes. The application was now fully functional and matched the specifications outlined in the design document (19:22).

While testing the app, I realized that viewing a single search result page was helpful, but having the ability to navigate to nearby pages would provide



better context. Since the full PDF was already loaded in the browser, adding this feature was a natural next step.

The assistant quickly added new tests and implemented navigation buttons (see the following screenshot), allowing users to move easily between pages (around minute 22).



Now we can bring this tool together with the CLI search system we built in the previous chapter. We still have a few small differences to iron out, but they fit together nicely—in fact, they already do in my own setup.

You can check out the complete code on the Git repository¹⁰ or try integrating everything yourself with the help of your AI assistant. Either way, it's a great exercise to reinforce what we've covered so far.

10. <https://github.com/uberto/lessui>

What We Covered

This chapter gave us a practical look at what it's like to build a real application with an AI assistant. Here's a quick summary of what went well and what still needs attention when working this way.

What Worked Well:

- We built a fully functional application in under an hour, even though we used a programming language and framework we weren't particularly experienced with.
- It was a great opportunity to get hands-on exposure to Elixir and Phoenix. The assistant helped us stay productive while learning, reducing the usual friction of ramping up on a new stack.
- The overall plan held up. The agent handled development tasks effectively, resolved issues, and self-corrected when prompted.
- The documentation generated by the assistant was clear, complete, and, honestly, better than what we would usually write by hand.

What Didn't Work So Well:

- We had to regularly repeat or reinforce development process instructions. The assistant often tried to skip tests or documentation in favor of generating code faster.
- The assistant had a tendency to assume things were working without verifying—skipping test runs, UI checks, or even proper commits.
- Occasionally, the assistant added unnecessary complexity (like mocks) or forgot earlier decisions, such as UI naming or reusing existing test files.

Despite these hiccups, the overall outcome was solid. With the right guidance and supervision, the LLM was a fast, capable coding partner. And with a clear process in place, we were able to stay in control and move quickly without compromising quality.

Perhaps most importantly, we proved that it's entirely possible to learn a new language and framework while shipping real features, as long as we combine curiosity with structure and let the assistant fill in the gaps.

In the next chapter, we'll shift our focus from greenfield projects to existing codebases. We'll explore how AI tools can help navigate large, complex, or legacy projects—from understanding unfamiliar code to refactoring, testing, and gradually modernizing systems without breaking them.

Working on a Large Codebase

After building a relatively complex application in the previous chapter, it's now time to focus on working effectively with large, existing codebases using AI assistants. You'll learn how to set up clear documentation and manage context carefully so the model doesn't get lost in the complexity and can provide useful help.

This chapter is all about practical techniques for everyday tasks: making safe code changes, refactoring small parts without breaking everything, and speeding up the tedious but important work. We'll look at how to write tests for tricky edge cases, migrate from one library to another, and introduce new abstractions to clean up duplication.

The goal isn't to find a magical prompt that fixes legacy code in one shot. Instead, we want to make our day-to-day work in messy, complex systems less painful and more productive—and maybe even more enjoyable.

But before we dive in, let's look closely at a couple of useful concepts that will guide us through this chapter.

Managing Technical Debt

Technical debt is what piles up after we've made fast decisions in the past, trading long-term quality for short-term speed. Those shortcuts were often necessary and helped us ship quickly, but over time they start slowing us down.

It's like having a debt with compounding interest. The longer we ignore it, the more it costs.

That's why managing debt isn't a luxury for teams with extra time. It's essential. It keeps the project healthy and helps the team stay productive in the long run.

AI can help here by cutting the time and effort needed for most maintenance work. It can point out duplication, inconsistent naming, or missing tests, and make small refactorings less painful. It can't replace human judgment—and sometimes it makes huge assessment mistakes—but when guided well, it can significantly speed up cleanup and redesign.

In the end, managing technical debt isn't about perfection. It's about keeping the system understandable and flexible enough to move forward. AI assistants are a great tool for achieving that goal, for those who know how to use them.

Understanding Legacy Code

Technical debt makes our code harder to read and change. But worse is that sometimes we have code running in production that nobody understands anymore. It's been there for years, the people who wrote it have moved on, and now everyone's a little afraid to touch it. That's what is usually called legacy code.

Given this picture, it's no surprise that legacy code has a bad reputation. We describe it as old, unmaintained, confusing—something nobody dares to touch. But before we throw it under the bus, let's look at what it tells us. It works. It's in production. It delivers value. If it didn't, it should have been deleted long ago.

A quote I like captures this perfectly:

Legacy code often differs from its suggested alternative by actually working and scaling.

—Bjarne Stroustrup

Legacy code is often running critical parts of a company's business, so it's important to keep it in good shape. Unfortunately, that's not easy. It's rarely documented properly. It's full of patterns that once made sense but no longer do after years of changes. It's often poorly tested, and adding new tests is tough because everything is so tightly coupled.

Working on legacy code feels like trying to fix a machine while it's running, with no manual and half the parts renamed.

With a good process, an AI assistant can make this work more manageable. It can help us understand the codebase, suggest small fixes, and make the original intent clearer. It can also generate documentation as we go—sometimes for the first time—saving hours of digging through unfamiliar logic. That lowers mental load and frees developers to focus on redesigning and improving the code itself.

In the end, maintaining legacy code doesn't have to be a dreaded task. With the right tools and mindset, it can become a challenging—even satisfying—part of our craft.

Enhancing Big Code Projects with AI

When working on large codebases—especially those meant to be maintained for years by teams of dozens or even hundreds of developers—we need a different mindset than when building small greenfield projects.

These systems have history, complexity, and often a fair share of technical debt. Understanding them takes time and effort, so we need to be very careful about how we use AI.

Alberto Brandolini once expressed that it is the developer's understanding, not the domain expert's knowledge that becomes production code,¹ an insight that fits perfectly here.

What we don't want is production code that reflects only the LLM's understanding. That would make it legacy code from day one—because nobody would truly understand it.

AI can help us deepen our understanding, generate ideas, and speed up some tasks, but we must not delegate responsibility. It's tempting to focus on specs and let the AI handle the rest, but in the end, it's the code that runs—not our documents.

This capability alone fully justifies the cost of AI assistants and the effort required to learn how to use them well. Our real value as developers lies in how well we understand both the domain and the code. No AI—at least until it becomes sentient—can replace that. Until then, we need to review every change carefully and make sure it matches our mental model of the system.

For this reason, in my experience, the single biggest benefit of using AI assistants is the ability to answer questions about the codebase—for example, locating the function that determines a certain behavior, understanding why it behaves that way, or figuring out which part of the code should be changed to fix a related bug. It's also extremely useful for explaining complex code or frameworks you're not familiar with.

Beyond that, many other useful applications exist: updating documentation, writing or adjusting tests, performing complex refactorings, or generating large amounts of boilerplate code.

1. <https://x.com/ziobrando/status/634668319006683136>

But when it comes to writing high-quality, sophisticated code that strictly follows an existing architecture or set of guidelines, LLMs are currently not very strong. I wouldn't say it's impossible to make them do it, but it's often faster and simpler to write that kind of code yourself—especially if you enjoy writing complex code.

Preference is also a consideration. Some people prefer supervising the AI. Others get frustrated by spending their day reviewing AI-generated code. Or maybe you're like me, and it depends on the day.

Avoiding the Gambling Mode

Even with all the benefits we've discussed so far, a real risk is that adopting AI can lead to a loss of productivity and become a source of frustration. In my experience, the main problem isn't the tools themselves but how our own minds react to them.

It's worth repeating that if you find yourself in a situation where the code almost works, but the assistant has already gone too far and generated too much or overly messy code, don't be afraid to throw it all away and restart cleanly from the last known good commit.

These situations are the biggest time black holes when working with AI. It's always tempting to think that just one more prompt will fix everything. This is especially true because the code usually looks fine. AI-generated code is often clean, well formatted, and convincing at first glance.

On top of that, the typical frequent pauses make it easy to lose context and start trusting the assistant more than you should. Psychologically, you're pulled into an addictive feedback loop. Every small prompt produces a large amount of code, often accompanied by reassuring explanations and praise. You feel super productive, even when you're not.

In a conversation with Alberto Brandolini, he pointed out that our brain is attracted to a kind of “gambling mode.” It's not a good mindset for coding, but is very effective at building addiction. It also explains why vibe coding is so hard to stop once you start. The instinct is to double down, but that's exactly what we should avoid.

In addition to this is the *sunk cost fallacy*. After investing time and effort, it feels painful to throw away code that almost works and admit that the path was wrong. But at that point, cutting your losses is often the only winning move.

I know this from experience. It's a hard-learned lesson.

To make this more concrete, let me describe a typical workflow from my day.

A Typical Daily Flow

As an example, let's say I need to fix a bug or add a small feature.

I usually start by reading the ticket and discussing it with the AI in chat mode to clarify details and make sure I fully understand the problem I'm about to fix. Sometimes this requires digging around for more information, but let's assume the requirements are clear enough.

I then ask the assistant for a detailed analysis: what is causing the problem, what we should do to fix it, and how we'll test that it's fixed.

This analysis is often long and can take 10 minutes or more, but that's fine. At that point, I'm not yet "committed" to the task. I can do other things without harming my focus, because I haven't dug into the code yet.

When the analysis is done, I read it carefully. I also reuse it as the starting point for the implementation chat. From that moment on, I try very hard to avoid long pauses.

That's why I keep the AI assistant on a short leash. I ask it to implement the plan step by step. If a step starts taking too long, I stop it and ask more specific questions instead.

The goal is twofold:

- First, to avoid the classic LLM "loop of death."
- Second, to avoid long pauses—more than a few minutes—that would pull me out of the concentration zone.

At the end of the day, I don't want to commit any code to production that I don't fully understand and that I can't honestly defend in a code review.

Speaking of code reviews, one particularly useful feature is the ability to easily recall past conversations. When feedback arrives on a pull request, we can go back to the original discussion that led to the change and continue it with the new comments. This keeps context intact and makes revisions much more deliberate and controlled.

Given the complexity of big projects, we won't walk through a full end-to-end example here. That would be hard to follow unless you've already worked on the same system. Also, different projects require different techniques, and no single example could realistically demonstrate them all.

A Note on Pauses

What I noticed is simple: pauses are usually bad for focus. But the worst case is switching to an activity that takes real mental space and makes me forget what I was doing.

Getting a coffee is fine. Starting to learn a new language is not.

For the same reason, I've found it unproductive to work on two different tasks in parallel with AI. Whatever time I "saved" during the waiting moments, I later lost trying to reload the context in my head.

What I've seen to be a terrible waste of time is keeping the LLM busy with more prompts while not following what it's doing, and doing something else in the meantime. The result is often a mess. In the end, I still have to study the code carefully—sometimes more than if I had stayed focused from the start.

So instead of one long session, this chapter will focus on a series of practical, focused recipes—techniques that are useful when working on large, mature codebases with the help of AI. So let's dive in.

Starting Clean

Before we get into the more interesting stuff, let's start with recipe zero. It might sound like a repeat of what we've already covered, but better safe than sorry, especially when you're working in a team with other people.

Don't start prompting until the code is in a clean, working state. If the project is half-edited, uncommitted, or broken in some corner, the assistant will almost certainly get confused. It might misunderstand the current state, suggest changes in the wrong places, or make everything worse.

So before doing anything else, save your work and commit whatever is working. That gives you a clear baseline and a safety net. If something goes wrong, you can always roll back.

Just as important, never commit code you haven't reviewed or aren't sure about. Even when the assistant produces something that looks fine, it still needs your judgment. One careless commit can create a chain of problems down the line.

And finally, be clear about what the assistant should touch. Tell it where to work and what to leave alone. The more precise your instructions, the better the result—and the less cleanup you'll need after.

Working in Small Steps

The first recipe is simple: when working on a large project, it's better to move in small steps. This is different from greenfield development, where we can often build things in bigger chunks. In legacy or complex systems, small, focused changes are safer and easier to manage.

At the current state of technology, we can't expect the AI assistant to understand how everything works in a large, tangled codebase. It won't magically guess the architecture or internal conventions, and even if we document them explicitly, the system might simply be too big for the model to hold in context.

That's why we need to guide the assistant precisely. Instead of asking for entire features, we ask for a specific refactor, a test for one function, or a small utility method. Keeping the scope tight leads to better suggestions and fewer surprises.

Another challenge is that we often don't know the codebase very well ourselves. That makes it easy to give the assistant impossible instructions—asking it to modify something that doesn't exist, or to write code that goes against the project's conventions. This is where the assistant's behavior becomes important. Rather than framing it as a “senior expert engineer who always knows best,” it's often more effective to tell the model something like this:

“You are a junior developer doing your best. If an instruction seems unclear, wrong, or unfamiliar, ask for more information before continuing.”

This kind of prompt encourages the model to be cautious and collaborative. It also reduces hallucinations, as noted in research from Anthropic,² where uncertainty and clarification are shown to improve output quality.

Even with clear instructions, though, LLMs still make small, beginner-level mistakes—especially in large or messy codebases. They might forget an import, leave in dead code after a refactor, or accidentally duplicate something. That's normal. In these cases, it's almost always faster to spot and fix the obvious issues yourself rather than trying to describe the problem in detail in natural language.

Another useful trick I often use is writing pseudocode directly into the file—just a few lines that explain what I want to happen. Then I ask the assistant to turn that into real code. This is often quicker and more precise than trying to explain everything in plain English inside the chat window.

2. <https://arxiv.org/pdf/2507.21509>

Of course, you can use small-step instructions in smaller projects too. But in those cases, the assistant usually has an easier time figuring things out on its own, so giving it larger tasks or asking for complete features lets you move faster. As models continue to improve, they'll get better at handling bigger and more complex projects, so this balance will likely shift over time. But for now, in the messy real-world codebases most of us deal with, breaking things down is still the key to getting good results.

Documenting Behavior Through Tests

Let's take *Kondor-Json* as an example.^a It's a library I wrote to handle JSON precisely in Kotlin, without relying on boilerplate DTOs just to bridge the gap between domain models and JSON schemas. It's built around functional programming principles, which you can read more about in my other book, *From Objects to Functions [Bar23]*.

In this kind of code, where correctness matters and logic can be subtle, it's easy to miss edge cases. One effective use of AI here is to ask it to generate tests for tricky inputs or scenarios we might have overlooked. The assistant doesn't need to understand the whole system, just the function and its expected behavior.

Now let's use the LLM to write a few more tests. This not only helps verify behavior but also documents how the feature is supposed to work. Well-written tests can serve as a form of living documentation, especially in parts of the code that are hard to explain in plain comments.

For this example (and some of the ones that follow), we're using JetBrains IntelliJ IDEA with Junie,^b JetBrains' LLM assistant. It works especially well when dealing with Java and Kotlin code. Watch the video at <https://youtu.be/kUAagcC6HKE>.

We start by asking Junie to check whether any test cases are missing:


Junie prepares a step-by-step plan and starts analyzing the existing tests. It quickly spots that the converters for *BigInteger* and *BigDecimal* aren't fully covered, and also notes the absence of tests for special *Double* values like *NaN* and *Infinity*.

It then writes the missing tests, runs them, detects a small mistake in one of the assertions, fixes it, and reruns the tests automatically. Our role here is simple: review the changes (which only involve new tests) and commit the result.

a. <https://github.com/uberto/kondor-json>

b. <https://www.jetbrains.com/junie/>

As you can see from the image below this tedious task has been completed in just over five minutes, something that would've taken much longer by hand. Good job, Junie!



Understanding with Ask Mode

LLMs are great for small well-defined changes, but what if we don't know the code well enough to even ask for specific tasks?

That's where the second recipe comes in. We can use AI to help us learn and get familiar with a new codebase, language, or library. This approach is especially useful when we've just joined a project or are working in unfamiliar territory.

The idea is simple: instead of asking the assistant to write code to solve a problem, we ask it to tell us how to do something without writing any code, and then we write it ourselves. In a way, it's the opposite of the previous recipe. But there's a big advantage: by typing the code ourselves, we absorb much more of what's going on around it. We understand the structure, spot conventions, and notice patterns we'd easily miss by just reading or copy-pasting.

Since we're the ones writing the code, we don't have to worry about hallucinations (hopefully!). And if something isn't clear or doesn't work, we can ask the assistant for clarification. It becomes a patient, helpful domain expert, rather than a code generator.

Most assistants even have an Ask mode, which answers questions without producing code and instead gives us a small document with a plan we can follow step by step.

Fixing a Tricky Bug in Kondor-Json

To put the “ask the AI” recipe into practice, I recorded a session where the AI helped me fix a tricky bug in the new version of Kondor—my JSON library for Kotlin, which we saw a few pages earlier. Watch the video at <https://youtu.be/HrZhIfHTeU>.

A bit of context will help: Kondor has a converter that specializes in translating a field with a map of key-value pairs into a raw JSON object, and vice versa, without needing to know in advance what the map contains.

In version 4.0, a major change was introduced: instead of first converting JSON into a tree of `JsonNodes` and then into Kotlin objects, the new parser goes directly from the raw JSON string to the target objects. This provides a big performance boost (the main reason for the 4.x rewrite), but it also required a lot of careful internal changes.

One of the trickiest parts was updating the `JMap` converter. I had forgotten some of its internals and was struggling to adapt it to the new architecture.

Enter the AI. Rather than asking it to “fix the code,” I asked Junie (JetBrains’ LLM assistant) to explain how the code needed to change. It offered a few suggestions, some of which were subtly incorrect. But even those half-right answers were helpful. Writing them down and thinking through the logic helped me see what was missing, and ultimately led me to the correct solution.

As you can see in the video, we tried several dead-end approaches before finally reaching a clean solution. That’s part of the process—and the learning. Not only did I fix the bug faster than I would have on my own, I also deepened my understanding of how to convert to the new parser in the process.

This approach is a bit like rubber duck debugging, except the duck talks back (even if it doesn’t always get everything right). So you can use it every time you want to not only fix the bug but also get familiar with the code and understand why the bug happened.

Enjoying Ping-Pong

The third recipe is a combination of the first two. When we have some understanding of the codebase but still don’t feel fully confident, we can switch back and forth between asking and coding, like a game of ping-pong.

We start by asking the AI to outline how to approach a specific problem—just like in Ask mode. Once we’re happy with the plan or direction, we switch modes and ask it to implement the idea in small, focused chunks.

This back-and-forth rhythm keeps us in control while still moving fast. It’s a flexible, pragmatic way to combine exploration and execution—building confidence without losing momentum.

Interestingly, this is also how some models work internally with their Planning mode: the AI first creates a plan to solve the problem, then follows that plan step by step. By putting ourselves in the loop, we get the same benefits while also learning more about an unfamiliar codebase.

Let's look at a concrete example.

Fixing the New Data Class Converter

While working on the new version of Kondor, I began developing a specialized converter for Kotlin data classes. The goal was to avoid explicitly calling constructors in every converter by using reflection to grab a reference to the primary constructor the first time it's needed, and then invoke it with arguments extracted from the JSON. (Watch the video.³)

Everything worked fine, except for one issue—nullable arguments. If a nullable field wasn't present in the input JSON, the constructor was invoked with the wrong number of arguments, leading to runtime errors.

To fix this, I first asked the AI for a detailed explanation of what was going wrong. Once I clearly understood the mechanics behind the error, I was able to write a precise description of what needed to change in the constructor-invoking logic. Then I asked Junie to apply the fix.

At that point, it worked beautifully. Earlier attempts where I simply asked the AI to “fix the bug” had failed. What made the difference was the ping-pong approach: understand the problem first, then guide the change with clear intent.

Another important aspect here is Kondor is an open source library. I'm not satisfied with it merely working. I want the code to be concise and elegant so it's easier to maintain and easier to extend without pain. AI is a great help in this process, especially when used with clear human judgment and direction.

This pattern works especially well in complex scenarios where the AI needs help to translate an analysis of the problem into focused steps to solve it. By actively collaborating instead of passively delegating, we make better use of the assistant's capabilities while keeping using the human brain where it excels.

Providing Feedback with Acceptance Tests

Throughout this book we've seen the same pattern repeatedly. Most of the time, the code generated by an assistant is very good. Every now and then,

3. <https://youtu.be/yyXxcj-NHMg>

though, it's completely off. This is what is often called AI slop, and it can be hard to spot by just skimming the code.

This problem becomes much more frequent in large or legacy codebases. The reason is simple. LLMs work on tokens. When analyzing and generating code, they can easily be sidetracked by similar-looking patterns or fragments that serve completely different purposes. Without feedback, the model has no reliable way to know it went in the wrong direction.

The most effective solution is to introduce feedback. We need to give the assistant a clear and objective signal about whether the application is behaving correctly so it can self-adjust. A good way to do this is through acceptance tests, where we verify that the overall behavior of the system is acceptable.

The easiest form of acceptance tests is end-to-end tests. They validate the system from the outside, in the same way a user or another system would. More importantly, they give the assistant something concrete to run and observe. Instead of asking us to manually check whether a feature works, the agent can write the test, run it, and inspect the result by itself.

In practice, this means letting the assistant create or extend an end-to-end test that exercises the feature we're working on. The agent runs the test and iterates on the code until it passes. Our role is then to review the test before committing it. We check that it covers the intended behavior and that the right assertions are in place.

This review step matters a lot. We need to pay close attention to what is asserted. AI assistants often write tests that technically pass but assert very little. This may reflect their training data, where many real-world test suites are shallow or incomplete. A passing test is only useful if it fails when behavior is wrong.

The main downside of relying heavily on end-to-end tests is speed. As the application grows, these tests tend to become slower and more expensive to run. This hurts feedback loops, both for humans and for AI agents. UI-based tests can also be hard to read and maintain, especially when they depend on fragile selectors or layout details.

Can we do better? Yes, at least if we're willing to invest in our architecture.

Instead of tightly coupling the application to the UI, we can often adopt a command-based approach. Rather than driving the system through clicks and screens, we expose a set of commands that represent meaningful user actions. These commands allow us to exercise the system much faster and in a more controlled way than full UI-driven end-to-end tests.

Command Query Responsibility Segregation

An architectural pattern that fits particularly well with AI-assisted development is CQRS, Command Query Responsibility Segregation.

CQRS separates commands, which change state, from queries, which only read data. This gives us a clearer and more explicit model of what the system can do. When commands are well named and focused on side effects, it becomes much easier for an assistant to interact with the system correctly.

I cover CQRS and related functional design patterns in more detail in my earlier book, *From Objects to Functions* [Bar23], where I explain how these ideas map naturally to functional-style design. The same principles that improve testability and clarity also make systems much more friendly to AI assistance.

This gives us the best of both worlds: we keep strong end-to-end or acceptance-level coverage while maintaining fast feedback loops. Most importantly, we give the AI assistant a clear and reliable signal about what we expect from the application, in a form that’s also easy for us to verify.

Planning Your Bug Investigations

A variation of the previous recipes can also help us approach debugging more effectively. In general, LLMs aren’t great at debugging on their own. They often pinpoint where the problem appears but then try to fix it right there—usually by adding some ad-hoc condition. As you can imagine, that rarely fixes the root cause. It breaks layers of abstraction and sometimes introduces new bugs elsewhere.

When the tests still fail, the assistant quickly loses focus. Instead of narrowing in on the root issue, it starts making random or speculative changes—rewriting nearby code, “fixing” things that weren’t broken, or even breaking working parts without making any real progress.

This happens because the model is trying to be helpful in a situation where the problem can’t be solved in that spot. LLMs don’t reason through why the failure happens—they lack a mental model of cause and effect. They optimize for “change something that looks relevant,” not for understanding the system as a whole.

To address this, we can shift the LLM into a different role—not as a code modifier, but as a planner. Specifically, we can ask it to create a step-by-step debugging plan that follows a progressive logic, starting with the most likely causes, and then eliminating them one by one.

Surprisingly, this works best without showing the model any code. That way, it won't "cheat" by hallucinating based on incomplete or irrelevant context. Instead, it's forced to "reason" purely from its knowledge and the high-level problem description.

Once we have the plan, we can follow it ourselves. After all, real-world debugging often involves more than just code; it requires looking at logs, testing in different environments, and sometimes even talking to users or stakeholders.

Let's walk through a real example from my day job. I've adjusted some details for privacy, and I can't show the code, but the core situation is exactly as it happened.

Fixing an Inexplicable Bug

We were dealing with a strange production bug: a business rule that applied a discount to returning users wasn't triggering, even though all the conditions looked right. The rule's logic seemed fine in the code, and the unit tests were passing. (Watch the video.⁴)

After a bit of head-banging on this one, I asked the LLM to help by generating a structured debugging plan. Here's the prompt I used:

"In production, a business rule for applying a discount to returning users did not trigger at the right date with the right user."
Please generate a progressive debugging plan by eliminating possible causes step by step.
Start with the most likely cause and go toward the least likely.
Include validations for each step and what to do based on the outcome.

The LLM suggested nine steps in progressive likely order:

1. Check the rule implementation and confirm correctness via unit tests.
2. Ensure the rule actually triggered in production.
3. Confirm runtime data provided to the rule was accurate.
4. Verify the correct code version was deployed and active.
5. Check for failures in databases, caches, or feature toggles.
6. Inspect issues with date/time boundaries or timezone conversions.
7. Confirm tests themselves aren't flawed or giving false positives.
8. Look for unusual data states specific to affected users.
9. Inspect rare infrastructure or transient runtime issues.

4. <https://youtu.be/Zo7HbTFd5LE>

I had already verified steps 1, 2, and 4. So I worked through the rest and found that the issue was caused by a temporary failure in a cron job—the last suggestion of the list, of course (it always is).

Sure, I might have found the bug eventually on my own, but this structured plan helped remove a lot of the frustration and probably saved me a fair chunk of time.

So keep this recipe in mind the next time something strange happens and you're running out of ideas. Having a clear checklist to fall back on can turn a frustrating mystery into a solvable problem. LLMs are great at generating comprehensive debugging plans—even pointing out things you might have overlooked.

Designing in Parallel

The advent of AI has unlocked new ways of coding that we're only beginning to understand. One of them, for me, is how to explore design ideas faster and more freely.

Here's how it works: instead of aiming for a single clean solution, you ask the assistant to write or rewrite a feature in a few different ways. You might explore variations in design, structure, or library use. You review what it gives you, tweak the request, roll things back, and try again. This process can go on for an hour or so, and by the end, you've seen several concrete, compiling variations of the same idea.

Using multiple assistants and git worktree (or just checking out the project multiple times), we can even work on several branches at the same time. Most of these versions won't be complete or ready for production—and that's fine.

The goal isn't to build something final, but to quickly explore what's possible. You're basically running a set of tiny AI-guided spikes that help you test ideas and find your design direction, without sinking hours into building full prototypes.

This technique is especially helpful when you're stuck or unsure about architecture choices. It lets you shift into exploration mode without losing momentum. And because the assistant handles the repetitive parts, you can focus your own energy on the decision-making.

Once you've identified the parts that require real care—maybe a tricky state transition or a performance-sensitive piece—you can step in and code those directly, or at least check them very carefully. Then, hand things back to the

assistant to help wire everything together, clean up the boilerplate, and shape the final version.

This kind of serendipity coding doesn't replace solid design thinking. But it does make it easier to reach better ideas faster, often with fewer dead ends.

One practical note: parallel designing can burn through tokens very quickly. Depending on your assistant plan, it can become expensive. It's best used intentionally, for the parts of the system where exploration really pays off.

Explaining Unknown Libraries

When working on real-world code, we often rely on libraries that the AI model doesn't recognize—either because they're too new, too old, or internal to the organization. In those cases, the assistant struggles to help. It simply doesn't know how to use them.

Fortunately, there's a straightforward solution: we need to feed the assistant some documentation. Ideally, this includes not just API references but actual usage examples. Giving the LLM concrete examples provides the context it needs to reason about the library properly. In fact, many modern libraries now include LLM-friendly documentation files meant specifically for this purpose.

In the most extreme case, we might be writing the library ourselves. Of course the model won't know about it—we haven't even finished it yet. But that's fine. It just means we need to treat our documentation not only as something for human readers but also as a way to teach our AI assistant how to help us more effectively.

Let's walk through a real example to see how this works in practice.

Using Kondor in a New Project

Going back to vibe coding, let's build a simple tool that lists Studio Ghibli movies using the public API.⁵ But this time, we'll use Kotlin with Kondor-Json for parsing the data.

Initially, I asked the LLM to generate some code using Kondor-Json—without giving it any extra context. The results were mixed: some parts were correct, others clearly hallucinated. This isn't surprising. Kondor-Json isn't one of the most widely-known libraries, so the model likely had only sparse or out-dated examples in its training data.

5. <https://ghibliapi.vercel.app>

But after I pointed the assistant to the online document⁶ explaining exactly how to use the library and asked the same question again, the model generated the correct code, and it ran perfectly—displaying a neat list of my favorite Japanese anime.

One new thing to keep an eye on is the proposed standard for LLM-friendly documentation.⁷

Many libraries, including my own, Kondor-Json, have already adopted it. With this approach, LLMs can access documentation in real time—always up-to-date, version-aware, and easy to query. That means no more outdated examples or version mismatches. It may even be the ultimate solution to prevent LLMs from relying too heavily on Stack Overflow.

Scaling Refactors with Scripts

One of the main limitations of LLMs when working on large projects is how little they can handle at once. They have a limited context window and can only take a few steps before losing track. For example, if you ask the assistant to fix something simple, like the order of imports in a file, it will likely do a good job. Maybe it can even handle a whole module with a dozen files. But if you ask it to apply the fix across a project with thousands of files, it might stop halfway through, or worse, tell you it's done when it actually touched only a few.

Now, fixing imports is the kind of thing our IDE can usually do faster, more reliably, and without help. But when it comes to more nuanced refactorings—things that require understanding the code—LLMs can still be valuable. The trick is to shift from working on individual files to generating scripts.

Instead of asking the assistant to refactor everything directly, we ask it to write a script that can iterate through all the files and apply the transformation one by one. This approach lets us scale the LLM's intelligence across the whole project while keeping things under our control.

Let's look at a real-world example.

Migrating from One Library to Another

http4k is a large open source Kotlin library for building web servers and clients in a functional style. It includes over 180 modules and a very large number of tests.

6. <https://github.com/uberto/kondor-json/blob/main/lms.txt>

7. <https://lms.txt.org/>

At some point, the maintainers switched their assertion library from Hamkrest to Strikt. So far so good, but thousands of existing tests still use the old assertion style. Migrating them is not conceptually difficult, but it's not trivial either. Hamkrest and Strikt have very different APIs and encourage different styles of assertions. A simple search-and-replace is not enough. At the same time, updating everything by hand would take days of focused work.

My suggestion to the maintainers is not to ask the AI to directly change the whole codebase but to ask the assistant to write a custom migration script. The script would parse test files, detect Hamkrest-style assertions, and rewrite them using the equivalent Strikt style as faithfully as possible.

First, we need a list of test files to process. This can be done by walking through the test directories of all modules and selecting files based on a regular expression. At this stage, we're only collecting filenames.

Once we have that list, we can process the files one by one. For each file, we send its content to an LLM with clear instructions on what needs to be changed. Using a command-line-driven assistant makes this easier to automate. For example, tools like Claude Code allow you to execute a prompt directly from the CLI. The script itself doesn't need to understand Kotlin deeply. Its main job is to call a templated prompt along the lines of "rewrite all Hamkrest assertions in this file using Strikt, preserving behavior."

The script calls the assistant for each file and writes the result back to disk. This process is relatively slow compared to a pure text-based transformation, but it doesn't require supervision. If something goes wrong, we can always inspect the diff or roll back the changes.

If you want to try this approach yourself, I recommend starting with a single module. This lets you validate the prompt, the script, and the results on a smaller scale. Once you're confident in the outcome, you can run it on the rest of the codebase and open a pull request to the http4k project.

I expect this kind of task will become much easier in the future as assistants get better at using IDE refactoring tools directly and become easier to script—no need for a separate app or service.

But even today, it's not a big deal. And it's much faster than doing the same refactor manually, file by file.

Empowering with MCP

Sometimes it would be useful to connect our AI assistant to company-specific tools and data—things like documentation stored in Notion or Confluence,

logs in Splunk or Kibana, or tasks and tickets in Jira or Asana. This kind of integration can unlock a lot of value.

For example, we could ask the assistant why a deploy failed. Instead of guessing, it can pull recent production logs, cross-reference them with open incidents, and suggest a fix based on an internal runbook. This moves the assistant from being a code helper to becoming a real part of the development workflow.

At first, this might sound like a feature reserved for future enterprise tools. But in reality, it's something we can already do—today—with just a bit of setup. That's where Model Context Protocol, or MCP,⁸ comes in.

MCP is a simple and powerful standard that lets us expose structured functionality to AI models through a small web server. It works a bit like a plugin system, but it's more flexible and easier to build for. The core idea is an MCP server advertises a list of capabilities (like “search documentation,” “get open incidents,” “query logs”), and each capability describes how it works, what parameters it expects, and what kind of result it returns.

Because MCP is self-describing, the model doesn't need to be pretrained on your specific company systems. Instead, the agent reads the list of capabilities directly from the server, understands what's available, and knows how to call each function just by reading the schema. The assistant effectively learns how to use your tools on the fly.

A useful way to think about this is that we're teaching the assistant new skills. We describe those skills in a standard format and expose them over HTTP. That's all. There's no need for complex SDKs or manual integration work.

This opens up a wide range of possibilities. We can build small services that expose internal documentation, operational data, incident status, or support tickets. Once a service is available through a well-formed MCP interface, any compatible LLM agent can start using it immediately.

Let's look at a small hands-on example.

Building an MCP Server

As an example, we can build a simple MCP server to give the assistant access to the documentation we stored in the LESS vector database.

8. <https://modelcontextprotocol.io>

To do this, we'll use the FastMCP Python library, which makes it quick to get something running. The MCP server will expose a single capability—something like `search_less_docs(query: string) -> result: string`. The FastMCP library handles the details of making this self-describing for the assistant.

MCP servers can expose their API either as an HTTP server or as a Stdio command-line application. This second way is much safer and is the best option for something local.

Once it's working, we can tell the assistant how to launch it, and the assistant can request the MCP manifest with all the details to the MCP server itself. Once done, we can ask questions like this one:

How do I access a database in Kotlin in a functional way, without side effects?

The assistant will query the documentation server behind the scenes and answer using real, context-aware info.

You can extend this same pattern to integrate with anything your company uses: ticketing systems, internal APIs, config databases, chat histories, dashboards—you name it. MCP makes it easy to build your own toolkit, one endpoint at a time. For example, you can ask your assistant how to fix a bug in production by analyzing the logs for a specific error. Or have it help implement a feature by referencing a ticket ID, pulling in the relevant context automatically, and even updating the ticket with useful information once the work is done.

You don't need to wait for vendors or platform support. You can build exactly what your team needs, starting today.

We can also connect this back to acceptance testing, as we discussed in [Providing Feedback with Acceptance Tests, on page 75](#). Instead of only running tests offline, we can allow the AI agent to run the application in real time through an MCP that abstracts over application commands. The agent interacts with the system programmatically, observes the results, and iterates quickly. Once a feature is complete and stable, we can store the sequence of commands as a new acceptance test.

That said, you don't need MCP for everything. Tasks that assistants already handle well—like accessing the file system, running shell commands, or using a SQL console—don't require extra endpoints.

It's also important to be selective. As Simon Willison points out,⁹ adding too many MCP tools at once can hurt performance. The assistant may struggle to decide which tool to use, leading to slower or less relevant results.

9. <https://simonwillison.net/2025/Aug/22/too-many-mcps/>

Generating Documentation and Diagrams

As a final recipe, let's not forget that LLMs can generate any kind of text, not just code. That means they can take care of a lot of the repetitive or tedious work that usually slows us down.

We can use them to create sample data for tests, generate translations, convert JSON into usable code, build API mocks, or write quick documentation snippets.

These aren't always glamorous tasks, but offloading them to the assistant saves time and mental effort. It helps us stay in flow, instead of constantly shifting focus to do grunt work. Whenever you catch yourself doing something mechanical or repetitive, it's worth asking: could the assistant do this for me?

As an example, let's look at Kondor for the last time.

Documenting Kondor Better

The documentation for Kondor was still a bit sparse, especially in the Outcome module: a lightweight version of Either for Kotlin, with support for chaining and error accumulation. Many of its methods had no documentation beyond the unit tests and surrounding code.

With an LLM, adding proper Kotlin Doc comments becomes almost effortless. (Watch the video.)¹⁰


We can go a bit further and ask for some docs to make everything easier to follow. For each module, AI wrote a short markdown file explaining what it's for, what it's in charge of, and how it links up with other parts of the system.

But maybe the most helpful part is it also created some Mermaid diagrams to show how the parser, tokenizer, and converters work together. Instead of trying to explain all the moving parts in words, the diagrams give you a quick visual of the flow: what calls what, how data moves around, and where each piece fits in. It really helps when you're trying to get your head around the system or explain it to someone new.

The generated documentation still needed editing: trimming duplicates, adding clarity, and rewording some explanations. But overall, it was fast, useful, and a massive improvement over nothing.

As you can see in the diagram on the next page, the results look pretty good.

10. <https://youtu.be/Cf-BSd9t6qc>



This is honestly one of my favorite uses of LLMs. I enjoy writing code, and even writing books, but I’ve never loved writing documentation. Now I can get it done in a breeze—that is, I can ask the AI to do it for me.

Is It All Worth It?

We have these 10 recipes, but you may be asking yourself: even if we apply them correctly, is there an advantage to using AI assistants extensively on a big codebase, considering that we still have to check all the code ourselves in the end?

The answer is “Yes, but ...” In most cases there’s a clear advantage, but there are exceptions. Here are the areas where AI assistants can boost productivity the most:

- *Explaining the codebase:* If we’re not familiar with the code—maybe we’re onboarding, or maybe nobody knows it anymore—an AI assistant can answer questions and explain how it works almost as well as a knowledgeable developer. Just finding the right place to make a fix can take hours; cutting that time to minutes is a huge win.
- *Fixing trivial tasks:* Day-to-day change requests are often straightforward, sometimes boring—“add a field,” “check one more condition,” “add an observability point.” An AI assistant can often handle the whole task from just the ticket description—tests included. Our role becomes reviewing the logic, refactoring if needed, and committing. It’s faster, and even if it takes the same time, we can work on something more interesting or important in parallel instead of being stuck on trivial changes.

- *Solving esoteric issues:* Sometimes a library doesn't work as expected, or we hit a nasty, hard-to-trace bug. With their encyclopedic knowledge, AI assistants can find answers far faster than a human digging through documentation or forums. In the worst cases, this can save days.
- *Writing and updating documentation:* This one's self-explanatory, but worth saying. Most developers aren't great at writing documentation, and technical writers are expensive. Having the ability to keep comprehensive documentation up-to-date with minimal effort is a big win for any project.

What's left out? Two areas stand out: high-level product decisions and difficult design changes. In these cases, AI assistants are hit and miss; they can give you an insight but they often suggest terrible ideas.

As Kent Beck (yes, him again) famously said, "For each desired change, make the change easy (warning: this may be hard), then make the easy change."

AI is very good at making the easy changes. What it still cannot do reliably is figure out how to make a hard change easy in the first place. That part still requires our judgment and experience.

A genuine concern is whether using AI makes us forget how to program. In some ways, yes. It's similar to getting used to driving with a navigation system. Over time, you stop memorizing routes because you rely on the navigator. But you don't forget how to drive. If you lose connectivity, you can still manage with a map and a few stops to ask for directions, even if it takes longer.

Programming with AI works in much the same way. We may forget some syntax details or lose muscle memory for boilerplate patterns. This shouldn't worry you. What matters is keeping a solid understanding of design, trade-offs, and system behavior. As long as that stays intact, the rest can be recovered when needed.

Finally, the next page contains a quick reference to help you keep all the techniques we've covered handy at a glance.

What We Covered

This chapter has been all about learning how to work better with large, existing codebases using AI. We looked at a series of small recipes. Each one is simple on its own, but together they're surprisingly powerful.

Cheat Sheet for *Process Over Magic*

Agent Rulesets: Do's and Don'ts for Harnessing AI

- **Don't Tell It Who It Should Pretend to Be:**
"You are an expert that writes code carefully and never makes mistakes" sounds confident but doesn't improve output quality.
- **Instead Give Clear and Actionable Rules to Follow:**
Never commit code unless all tests pass.
Refactor only implementation; never modify test assertions.
- **Don't Ask It to Think:**
"You must proceed autonomously and complete the task correctly" sounds nice, but it won't change the LLM's limitations.
- **Instead Tell It to Double-Check when Unsure:**
"If the request is ambiguous, ask for clarification before proceeding. Explain why something won't work instead of guessing."
- **Don't Rely on Its Knowledge:**
"Write professional-level code" doesn't say much; everybody has different opinions.
- **Instead, Specify the Coding Style and Technologies:**
"Use functional Kotlin style with sealed classes and immutability, and http4k as web server."

Vibe Coding /Spec-Driven Development for POC

- Vibe coding using AI without checking the code is great for:
 - Proofs of concept.
 - Trying out new technologies or APIs.
 - Command-line utilities.
 - Internal tools.
- A more controlled approach is Spec-Driven Development:
 - Start with a clear, detailed specification.
 - Let the AI assistant generate code freely.
 - Validate the app running it (don't read the code).
 - Once working, commit it and repeat.

Greenfield Controlled Flow /One Prompt, One Commit

- **Only Commit Working Code**
Code must build, run, and pass all tests (unit +end-to-end) before committing.
If it fails, fix first—no adding features or scope creep until it's resolved.
Conversations with AI should stay focused on fixing the issue, not expanding.
- **Always Review Before Committing**
Carefully review all code changes before every commit.
Only commit code you're 100% confident in—ask the assistant for clarification if needed.
- **Focused Prompts**
Keep each prompt and commit focused on one task: either add a feature or refactor.
Avoid mixing goals—it leads to messy diffs and hard-to-revert changes.
Minor improvements? Note them for the next cycle.
- **When Stuck, Roll Back**
If things get messy or the assistant is stuck, roll back to the last clean state.
Try again with a smaller, clearer prompt to avoid cascading errors.

Recipes for Working with Large Code Bases

- **Working in Small Steps**
In complex or legacy code, make small, focused changes. This is easier to review and, if necessary, to roll back.
- **Understanding with Ask Mode**
Ask the assistant how to do something, then write the code yourself. This improves learning and keeps you aware of the surrounding code.
- **Enjoying Ping-Pong**
Alternate between asking for advice and having the assistant implement small pieces. Great when you're semi-confident in the codebase.
- **Providing Feedback with Acceptance Tests**
Let the assistant validate its own work using end-to-end acceptance tests. A clear command-query separation makes this much easier.
- **Planning Your Bug Investigations**
Ask the assistant for a step-by-step debugging plan, starting from most likely causes and narrowing down logically.
- **Designing in Parallel**
Ask for multiple variations of a solution to explore design options. Compare, tweak, and experiment to find the best fit.
- **Explaining Unknown Libraries**
Help the assistant by pasting docs and usage examples for unknown or internal libraries. See also: 1mstxt.org.
- **Scaling Refactors with Scripts**
For big changes, ask the assistant to write scripts that process files incrementally—works around context limits.
- **Empowering with MCP**
Use the Model Capability Protocol (MCP) to connect the assistant to internal tools like Notion, Jira, or logs—adds massive real-world value.
- **Generating Docs and Diagrams**
Leverage the assistant to write documentation, generate sample data, create diagrams (e.g., with Mermaid), and more—not just code.

Here's a review of what we covered:

AI is a helper, not a replacement

AI assistants aren't here to replace us, but to support us. They help carry the mental load, speed up the repetitive parts, and make it easier to navigate complexity without losing our grip on quality.

Fluency means knowing when not to use it

Becoming fluent with AI means also learning when it's better to hand off a task, and when it's simply faster and more effective to do it ourselves. Recognizing that boundary—when to delegate and when to take the wheel—is a skill that grows over time, and it's key to using these tools well.

A good process makes everything smoother

We saw how starting with clean commits and clear goals makes the assistant far more effective, and how working in small, focused steps helps us stay in control.

Use AI as a lens to explore the code

We can learn unfamiliar code and tools by asking, trying, and iterating, using the assistant to understand existing code, not just to generate new code. And when we're stuck, we can use it to explore alternative designs or quickly test new ideas, without wasting hours on dead ends.

The real gain is in time, not lines of code

Whether we're introducing new abstractions, migrating libraries, or just trying to understand a messy class, the assistant becomes a kind of power multiplier. It doesn't need to be perfect. It just needs to be fast and consistent enough to make our job easier—and most of the time, it is.

In the end, in big projects, the real power of AI isn't about how fast it can generate code. It's about cutting the boring parts, learning faster, and reshaping existing code more effectively. And when we can do all that so easily, is “legacy” a correct term anymore?

And that brings us to the next chapter: adopting AI at the team level. We'll see what it means when an entire team starts working this way and how we can support that transition.

Collaborating Inside an AI-Powered Team

In this chapter, we won't focus on the code generation. Instead, we'll look at how AI assistants are changing the way we work as teams and how we can adapt to make the most of it. This shift isn't just limited to tools—it touches team dynamics, workflows, and even how we think about roles.

It's still too early to fully grasp the scope of this transformation, but here are some thoughts on the changes that are already taking shape and how we might navigate them.

Adopting AI

When it comes to bringing AI into a team, we should support it but not force it. Not everyone is ready to accept change at the same time, and that's fine. What helps is creating space for people to explore and get comfortable. Organizing internal hackathons, setting up quick workshops, and sharing small wins can go a long way. The goal is to act as facilitators, not to enforce rules from the top. This is still a new field for most of us.

It's also important to have a shared set of guidelines. Nothing too fancy—just a clear Markdown file outlining how we use LLMs and agents, what's acceptable, and how that fits with the rest of the team's coding standards. When everyone refers to the same base document, it reduces confusion.

At the same time, we should acknowledge that each developer might work differently with AI. Some may prefer their own ways of prompting or reviewing suggestions. To support that, we can add a line to the shared rules file, like “If you use personal preferences or workflows, refer to `personal_rules.md` for more detail.” This gives space for customization without fragmenting the core workflow. Just make sure `personal_rules.md` is in the `.gitignore`, so every dev can have their version without creating noise in the repo.

The goal here isn't uniformity. It's creating a shared baseline that helps teams move fast without stepping on each other's toes. We want AI to be a boost, not a burden.

A good starting point is making it easy for everyone to try, tweak, and build trust in the process. Still, there's a limit to what organic adoption can achieve. Giving everyone a fair chance to learn the basics in a structured way ultimately benefits the whole team.

Training for New Tools

Learning to use an AI assistant well isn't especially difficult, but it does require time, experimentation, and a willingness to adjust old habits. Not every developer naturally leans into this, which is why dedicated training can be a real benefit for a team.

I've taught similar courses in several companies and at conferences, and I've noticed a few recurring patterns.

Some developers treat the AI with a kind of reverence. They see its output as something close to a calculator—precise, authoritative, and never wrong. Hands-on experience quickly changes this view. Once they see an LLM hallucinate, contradict itself, or make a silly mistake, they develop a healthier attitude: respectful of the tool, but not intimidated by it.

Others approach the assistant with deep skepticism—sometimes even hostility. They focus only on the moments when the model says something wrong or behaves clumsily, and dismiss the rest. Training helps these developers rebalance their perspective. Yes, LLMs can be stupid at times, but they can also save hours of work. If we use them constructively, they're a net win.

Another advantage is that structured training spreads practical techniques. Developers don't have to stumble through the same mistakes one by one. A short course can introduce the patterns, rules, and workflows we've covered in this book—and help teams develop their own process tailored to company needs. This cuts the time it takes for the whole team to become genuinely proficient with AI assistants of all kinds.

Moreover, the transition to an AI-assisted process brings challenges for both senior and junior developers—just in different ways.

Senior developers may feel the friction of unlearning habits. They've spent years typing everything themselves, relying on muscle memory and personal libraries of tricks. Switching to a more collaborative, prompt-based workflow can feel unnatural at first, and even slower, until they adapt.

Juniors, on the other hand, may struggle because they're still building their hands-on experience. That lack of confidence can become a limitation when reviewing generated code or deciding whether something “feels right.” They may just trust the AI.

That's why training can be so helpful: it gives everyone a safe place to practice how to collaborate with an assistant—how to give clear instructions, review output critically, and keep ownership of the code at all times.

The aim isn't to boost productivity with AI magic. It's to teach the process and mindset that help teams move beyond vibe coding and use these tools reliably, every day, avoiding all the traps and issues that we've discussed in this book.

Team Organization

If we look at typical software teams, the speed of writing code is rarely the limiting factor. Delays come from unclear specs, context switching, waiting on decisions, slow reviews, and manual processes in testing and deployment.

For this reason, adopting LLMs in software development isn't just about making developers faster; it necessarily affects all roles. Everyone needs to participate in the change—not only engineers, but also QA, designers, product owners, and analysts.

Otherwise, when coding gets much faster, it won't speed up the whole team. It will simply highlight everything else that's in the way. The theory of constraints makes this clear: improving something that's not the bottleneck can slow the entire system down, because the real bottleneck becomes even more congested.

Interestingly, this is reflected in real data. Several studies—including the DORA reports¹—have found no significant improvement in overall team productivity after adopting AI tools. Having experienced the benefits myself, I suspect the issue is not the technology but the lack of training and the lack of a broader rethink of how the whole team works.

So don't assume AI will automatically make your team faster. It can, but it requires effort, discipline, and changes across the entire workflow—not just faster coding.

The good news is that AI can help across all roles. Testers can generate automated test suites more quickly. Product owners and analysts can browse

1. <https://cloud.google.com/resources/content/2025-dora-ai-assisted-software-development-report>

and produce documentation with far less friction. Designers can prototype screens using AI tools. The entire team can benefit—as long as workflows evolve together.

It also makes sense to integrate AI into the tools we already use. It's now straightforward to have an internal chat that creates tickets from discussions or links directly to the relevant issue or commit. Many tools are already moving in this direction, and more changes are certainly coming.

And none of this requires anything close to Artificial General Intelligence. Even with today's limitations, there's plenty to improve. Automatic code reviews, for example, often flag incorrect issues—but even if only a third of the warnings are useful, that's still a big win for a very small cost.

Another example is documentation: keeping it up-to-date used to be considered impossible, but LLMs can update it automatically, and vector databases finally make searching in unstructured text both accurate and useful.

And, of course, building internal tools is now dramatically easier. With a bit of vibe coding, teams can create small, targeted tools that fit their workflow perfectly.

Hiring

Hiring is also going to change. LLMs are superhuman at solving algorithmic problems—the kind often used in traditional interviews. This not only makes it easy to cheat during remote interviews but also leads to awkward, even humiliating attempts to “prove” that no AI is being used.

In my experience, those kinds of interviews never reflected a candidate's real skills on the job. But now, they're not just flawed, they're obsolete. The same goes for trivia questions like “How many ways can you use the static keyword in Java?” That kind of knowledge is now instantly available and largely irrelevant.

We need interview processes that can't be gamed by LLMs. Instead of asking people to write code in a vacuum, we should assess how they work with AI. Can they break down a feature into clear, manageable steps? Can they review and improve generated code, spotting issues when they're present and knowing when to push back or try a different approach?

These are the new core skills we should look for, not memorizing obscure syntax or solving puzzles on a whiteboard.

For example, we could ask candidates to complete a task using an AI assistant, or review AI-generated code and spot issues or weaknesses. This gives us a clearer picture of how they'll perform in a real-world, AI-augmented environment.

And since LLMs can instantly fill in many technical gaps, traditional filters like “years of experience with X” start to matter a lot less. What we should be looking for are people who can think holistically about the product, communicate well, and make good technical decisions in context.

If you’re a candidate, AI can be a great training partner—it can help you learn a new technology for a take-home assignment, or act as an extra pair of eyes before you submit your code. But be aware of a risk: you can’t afford to be sloppy about understanding the problem or reading the code yourself. If you rely on AI for everything, what’s left that demonstrates your value? After all, the company can ask the AI too.

Onboarding and Code Review

One of the clearest wins from AI assistants is in onboarding of new developers. With good documentation, especially the kind the LLM can read and use, a new developer can feel like they have an expert sitting beside them from day one.

They can ask questions about architecture, modules, naming conventions, or coding patterns and get quick, useful answers. This doesn’t replace mentoring, but it fills the gaps and accelerates learning.

The same goes for code reviews. While we still need human judgment for overall design and verification, an LLM can act like a super-linter—spotting inconsistencies, catching stylistic issues, and reminding us of project-specific conventions. It’s like having a further pair of eyes focused on coherence and quality, without slowing down the process.

These tools won’t replace good documentation or strong team habits, but they can help scale them, especially in fast-moving teams or larger organizations.

AI and Developer Well-Being

Adopting AI tools brings clear productivity benefits, but it also introduces new sources of stress. Some are already well known, and others will likely emerge as we continue to adapt. Here are a few considerations:

- *Job security:* One common concern is fear of being replaced. People need to feel safe using these tools, not worried that by being more productive, they’re making themselves redundant. Teams need to be clear: AI is here to support developers, not eliminate them. The value is how we understand the system we build, not how many lines we type.

- *Flow and context-switching:* Some developers thrive on deep focus and long, uninterrupted blocks of work. Prompting and reviewing AI output often introduces small, frequent pauses, breaking that flow. For some, this can be frustrating or even exhausting. We need to allow room for different working styles.
- *Getting stuck in loops:* Ideally, AI should speed things up. But if we're not careful, it can lead to unproductive loops: generating and regenerating code, chasing small changes, and losing hours without real progress. This often happens without proper training on how to use AI assistants.
- *Brain shutdown:* Vibe coding—just accepting AI suggestions and going with the flow—can feel fast and even fun. But if we stop thinking critically, the quality of the code drops quickly, as we discussed in the first chapter.
- *Too much, too fast:* When the assistant produces hundreds of lines of code in minutes, it can feel overwhelming to review everything properly. That's normal. It's okay to slow down, take breaks, and trust that even at a calmer pace, you're still moving faster than before.

AI can be a powerful tool, but it's not a magic wand. We need to learn how to work with it in a healthy, sustainable way—one that keeps both our projects and our brains in good shape.

Security Risks

Using AI assistants inside a company brings clear productivity gains, but it also introduces serious security and privacy risks. These need to be understood and managed, especially when dealing with sensitive codebases, customer data, or internal tools.

Probably the first widely known example came from Samsung in 2023, when engineers accidentally uploaded proprietary source code and chip designs to ChatGPT while debugging. This wasn't an attack—just people trying to get their work done. The incident was serious enough that Samsung temporarily banned public AI assistants inside the company.

Since then, major LLM providers have introduced enterprise agreements with strict IP-protection guarantees, but the underlying risk remains: accidental disclosure happens easily when AI tools are part of the daily workflow.

These disclosure incidents fall into five major risk areas that every team should be aware of:

- *Unintentional data exposure:* This used to be mostly a developer problem, but not anymore. The most striking example came from Canada,² where clinicians were using ChatGPT to summarize patient notes and translate clinical documents, unintentionally sending protected health information to foreign servers.
- *Prompt injections:* Prompt injection has evolved into a multi-vector attack, blending traditional exploits with AI behavior. One 2024 demonstration showed how a model could be tricked into decoding a Base64-encoded script and inserting the malicious JavaScript into its own output—slipping past standard XSS filters.³ No jailbreaks, no trick questions—just a prompt engineered to do damage.
- *Model training leaks:* Unless you have explicit guarantees, data sent to a cloud LLM might be used for model improvement. The GitHub Copilot controversy—where it reproduced verbatim snippets of public repositories—showed that training data can resurface. Companies rightly worry that proprietary code might leak in similar ways. A 2025 analysis of Common Crawl—used by many commercial LLMs—found thousands of live API keys and passwords, including AWS credentials.⁴
- *Platform-level exposure:* Even careful users can be affected by bugs. Even if you use AI responsibly, platforms themselves can leak data. In August 2025, a configuration error caused private LLM conversations to become publicly accessible to search engines and archiving crawlers—without the users’ knowledge.⁵
- *Subtle vulnerabilities in generated code:* LLMs generate code quickly, but not always safely. A 2025 study across more than 100 models found that only 55 percent of AI-generated code was secure.⁶ The dangerous part isn’t just the vulnerabilities—it’s how confident and “correct-looking” the insecure code often appears. Worse, developers tend to trust AI-written code more, making review less thorough.

2. <https://www.insurancejournal.com/news/international/2025/11/21/848596.htm>

3. <https://blog.lastpass.com/posts/prompt-injection>

4. <https://www.computing.co.uk/news/2025/security/ai-training-dataset-leaks-api-keys-and-passwords>

5. <https://thefuturesociety.org/us-ai-incident-response/>

6. <https://www.veracode.com/blog/ai-generated-code-security-risks/>

The Lethal Trifecta

Simon Willison describes what he calls the lethal trifecta for AI agents:

- *private data*,
- *untrusted content*, and
- *external communication*.

When these three come together, the risks multiply. If an agent has access to your private data, processes untrusted content (for example, from users or the web), and can communicate externally—even something as simple as sending a request or an email—an attacker can trick it into leaking confidential information.

It's not the individual capability that's dangerous, but their combination. Understanding this pattern helps us design safer systems and keep AI tools inside clear, controlled boundaries.

To reduce these risks, we must not drop our usual security habits. On the contrary, we should be even stricter with things like these:

- *Avoiding sending sensitive data in prompts:* Never use customers' or any other sensitive data in prompts. Use placeholders or fake data whenever they're needed.
- *Knowing your LLM provider's data policy:* Always check whether prompts are stored, logged, or used for training. Choose providers that offer data control options, or consider on-prem solutions for critical cases.
- *Reviewing AI-generated code carefully:* Never assume the assistant got it right, especially for anything involving input handling, access control, or external integration. Review with a security-first mindset.

Finally, consider malicious attacks specifically targeting AI-generated code. They're already being used as attack vectors in new and unexpected ways.

For example, attackers have started publishing malicious libraries using names that are commonly hallucinated by LLMs—for instance, a fake package with a name like `http_helpers` or `string_utils`, hosted on public repositories like npm, PyPI, or Maven Central. These packages just sit there, waiting to be picked up by vibe-coded applications that trust the assistant blindly.

It's a new frontier, and the threats are real. Working with AI assistants doesn't just speed up development; it also expands the attack surface. So security needs to stay part of the conversation from the very beginning.

What We Covered

This chapter looked at how LLMs are changing more than just the way we write code. They're forcing a shift in how teams work, how developers learn, and how we think about security, tools, and roles. Here's a quick review of what we covered:

- *Speed shifts the bottleneck:* Faster coding doesn't speed up the whole team unless the rest of the process keeps up. To truly benefit from AI-assisted development, teams need to rethink the entire delivery pipeline.
- *Roles are evolving:* With LLMs leveling the field, years of experience matters less. The ability to guide the assistant effectively is becoming a key skill to evaluate.
- *Training and hiring need to adapt:* Both experienced and new developers must unlearn old habits and build new ones. Traditional interviews based on trivia or algorithm puzzles are now outdated.
- *Well-being needs attention:* The shift in pace and workflow can be stressful. Developers may worry about job security, lose focus from constant prompting, or feel overwhelmed by the fast output.
- *Security is a big deal:* LLMs introduce new risks and open up fresh attack surfaces. Staying secure requires awareness, good practices, and clear policies.

In short, LLMs are reshaping how we build software—not just faster, but differently. The challenge now is to evolve our teams, practices, and tools so we can take full advantage of what these assistants offer, without losing control or clarity along the way.

Demystifying Large Language Models

For better or worse, large language models (LLMs) have become almost synonymous with AI, even though they're the newest arrivals in the field. Among all technologies built on neural networks, they are the ones that have changed our daily lives the most.

In this appendix we'll take a look inside a large language model and see, at a high level, how it works. We'll skip the math and focus on the big picture. The goal is to strip away the aura of magic that surrounds LLMs and give us a practical mental model of what's going on under the hood.

We can absolutely use these AI assistants without knowing anything about their internals. But understanding the basics gives us an advantage. When we know how they “think,” we can guide them better and get more reliable results. Think of it as developing a kind of mechanical sympathy for LLMs.

Jackie Stewart, the legendary Formula 1 driver, used the term “mechanical sympathy” to describe a great driver's relationship with their car. You don't have to be an engineer to drive fast, but you do need to understand how a car works to be truly great.

In the same way, we don't need every technical detail about LLMs. But if we understand their basic principles, strengths, and limits, we can work with them instead of against them and get much better results.

Attention Is All You Need

Since the landmark 2017 paper “Attention Is All You Need,”¹ natural language processing using large models has exploded. Initially it was driven by the relentless scaling of model sizes and training datasets. But now, the trend is

1. <https://arxiv.org/abs/1706.03762>

shifting: instead of just making models bigger, research is focusing on refining training techniques and distilling models, making them smaller, more efficient, and less resource-intensive, and the results are very promising.

That said, improvements are becoming more incremental. Meanwhile, the number of players and tools in the AI space has skyrocketed. This signals that we're entering a phase of stabilization and commoditization. Running and developing LLMs is no longer limited to a handful of deep-pocketed giants—it's becoming accessible to smaller companies, startups, and even individual developers.

Another game-changing development is the rise of open source (or partially open) alternatives that are catching up fast with proprietary best-in-class models. This shift gives developers more control, customization, and cost-effective ways to integrate AI into their workflows. The AI landscape is becoming more democratized, and that's an exciting prospect for the future of software development.

What Are Large Language Models?

An LLM is a type of neural network designed to generate human-like text. It's built primarily on the Transformer architecture (by the way, GPT stands for Generative Pretrained Transformer), which enables it to process and predict words based on their context. These models are trained on vast amounts of text, code, and other digital content, encoding patterns, structures, and relationships between words into matrices of floating points. Unlike traditional rule-based systems, LLMs don't "understand" language in a human way—they predict the most likely next words in a sequence based on a given prompt and then choose one, often the most probable but not always.

The prompt is the initial input given to the model—what you type or say to get a response. The context includes both the prompt and everything the model has "seen" in the current interaction, including previous messages and system-level instructions. Often, LLMs also include a system prompt—a hidden message that guides the assistant's tone, behavior, or goals. In the case of AI assistants, the context can also include source files, documentation, or other reference material provided to help the model to "understand" the code.

On one hand, LLMs are impressive for their capabilities; on the other, they have incredibly frustrating limitations. They lack true understanding, reasoning, and intent, meaning they always generate plausible responses but have no way of knowing whether they're correct or not.

Wrong responses are often called hallucinations, but in reality, every response they generate is a kind of hallucination; it's just that we can recognize some as correct and others as incorrect, based on our understanding of reality.


The quality of their training data has improved dramatically in recent years, making them correct most of the time. Ironically, this makes their hallucinations even more dangerous because they're rarer and thus more difficult to detect.

The training and structure of LLM transformers is a fascinating topic, but we won't cover it in detail here. Still, it's useful to understand what an LLM ultimately is. In other words, what do you actually get when you download a model?

You can't download commercial models like ChatGPT or Claude, but many freely available models are on sites like ollama.com or huggingface.com.

Llama 3.3 Details

Let's take a look at a popular model, Llama 3.3, released by Meta.²



In the description, it says “70B parameters” and “Q4_K quantization.” This means the model contains 70 billion floating point values, each encoded using 4 bits (Q4_K is a technique for representing floating points in 4-bit precision). Based on this, the total size should be around 35 gigabytes (8 bits for byte).

However, the actual download size is 43GB. The difference arises because, as we will see, the parameters are grouped into two-dimensional arrays (matrices). While most matrices use 4-bit parameters, some use 6-bit, and a few use 32-bit parameters.

2. <https://ollama.com/library/llama3.3>

Now, let's see how an LLM uses these matrices to generate text. Understanding this will help us get better results when working with them.

When we give an LLM a question, the first step is splitting our text into tokens—small chunks, often one or two per word. Each token is then mapped to a list of numbers called an embedding. You can think of an embedding as the token's “coordinates” in a huge, invisible space. Tokens with similar meanings are close together in that space, even if they don't look alike as words.

For reference, Llama 3.3 has a vocabulary of 128,256 tokens. Each token is represented by 8,192 coordinates (dimensions). In other words, it encodes the meaning of each “word” in a precise 8,192-number fingerprint.

From here, the model starts looking at each token in the context and comparing it to all the others. This is the attention mechanism. It's like being in a meeting and deciding, for each sentence, which other sentences are important to keep in mind before you respond. The model does this for every token, at every step, constantly adjusting what it “pays attention to” as it builds the response.

For performance reasons, this attention happens in parallel across many heads. Each head focuses on a different part of the vector. In Llama 3.3, there are 64 heads, each one processing 128 elements of the vector.

After attention, the model applies a series of mathematical transformations (projections and other operations) to produce new matrices. This whole process is repeated layer after layer—Llama 3.3 has 80 of them—with each layer refining the understanding of the context and narrowing down what the next token should be.

In the final layer, the model produces a list of probabilities for every possible next token. Then it picks one.

If it always picked the single most probable token, the output would quickly become dull and repetitive. To avoid this, LLMs introduce randomness, controlled by a parameter called temperature. A temperature of 0 means fully deterministic output—always the most likely choice. A temperature of 1 mixes in more randomness, creating varied and more creative text. Each model has its own sweet spot for temperature, but we can adjust it depending on what we want (see the following image).

Name	Type	Shape
token_embd.weight	Q4_K	[8192, 128256]
blk.0		
blk.0.attn_k.weight	Q4_K	[8192, 1024]
blk.0.attn_norm.weight	F32	[8192]
blk.0.attn_output.weight	Q4_K	[8192, 8192]
blk.0.attn_q.weight	Q4_K	[8192, 8192]
blk.0.attn_v.weight	Q6_K	[8192, 1024]
blk.0.ffn_down.weight	Q6_K	[28672, 8192]
blk.0.ffn_gate.weight	Q4_K	[8192, 28672]
blk.0.ffn_norm.weight	F32	[8192]
blk.0.ffn_up.weight	Q4_K	[8192, 28672]
blk.1		
blk.1.attn_k.weight	Q4_K	[8192, 1024]
blk.1.attn_norm.weight	F32	[8192]
blk.1.attn_output.weight	Q4_K	[8192, 8192]

Magic Numbers?

Now, you might wonder: why 80 layers and 8,192 embedding dimensions? Why not 1,000 layers and 100,000 dimensions—or even more? The truth is, nobody really knows. On one hand, a higher number of parameters can make an LLM appear “smarter.” On the other hand, increasing parameters exponentially raises the cost of training.

We saw that jump from GPT-2 (with 1.5 billion parameters) to GPT-3 (175 billion parameters)—roughly a 100× increase. Then came GPT-4, estimated at around 1.5 trillion parameters—about 10× bigger again. After more than two years, GPT-5 arrived with an estimated 3–5 trillion parameters—only about 2–3× larger than GPT-4.

Clearly, there’s a lot we can learn about how to train these models more efficiently. When we compare them to the human brain—which operates with vastly lower energy consumption—it’s clear that we’re still far from optimizing AI in terms of both power efficiency and learning mechanisms.

Vector Databases

Let's quickly look at another interesting technology that's becoming useful with LLMs and natural language tasks: vector databases.

At the core, a vector database is just a database optimized for storing and searching vectors of floating-point numbers—that's where the name comes from. Many of them can also store strings and other metadata alongside the vectors.

We already saw that in a transformer neural network, each token is associated with an embedding vector that captures some aspect of its meaning. It's not just tokens, either—phrases, sentences, and even whole paragraphs can be turned into vectors too. That's basically how attention heads work under the hood.

Vector databases let us store these embeddings and then search for vectors that are “close” to a given query vector. Different algorithms exist for measuring this closeness, but the basic idea is always the same: compare how near two vectors are in their mathematical space.

Now, you might be wondering: why bother storing vectors at all? Here's why. Imagine you have a big chunk of text—a book, a set of documents, whatever. You can split it into smaller pieces, compute an embedding for each piece, and save all of them into the database.

Later, when you have a question or a search term, you don't have to scan the whole text. You just query the database for the closest embeddings. This gives you the parts of the text most related to your query, fast. You can either use those results directly, or pass them to an LLM to generate an answer—which is exactly what Retrieval-Augmented Generation (RAG) is all about.

Considerations

So, what can we learn from this about LLMs?

First, context size is crucial. Llama 3.3 doesn't remember anything beyond 131,072 tokens. This differs from human memory, which fades gradually while retaining important details. LLMs need constant reminders of relevant information.

Second, they generate text based on both their training data and the provided context—which, initially, is just our prompt. A vague prompt like “Write an application to search inside PDFs” will mostly rely on pretraining, meaning

the output is likely a slightly modified copy of an existing GitHub project with a similar description.

Finally, LLMs excel at recognizing and generalizing patterns, which is why they're so powerful. This is also why they can translate languages so well—preserving meaning while adjusting grammar and vocabulary. However, this same ability makes them prone to hallucinations, like assuming a Ruby library exists in Java by adding a “4j” suffix.

Unless a groundbreaking new algorithm is discovered, these considerations should hold true for future models as well.

Bibliography

- [Bar23] Uberto Barbini. *From Objects to Functions*. The Pragmatic Bookshelf, Dallas, TX, 2023.
- [FP09] Steve Freeman and Nat Pryce. *Growing Object-Oriented Software, Guided by Tests*. Addison-Wesley Longman, Boston, MA, 2009.
- [Gre25] Ben Greenberg. *Vector Search with JavaScript*. The Pragmatic Bookshelf, Dallas, TX, 2025.

Thank you!

We hope you enjoyed this book and that you're already thinking about what you want to learn next. To help make that decision easier, we're offering you this gift.

Head on over to <https://pragprog.com> right now, and use the coupon code BUYANOTHER2026 to save 30% on your next ebook. This offer does not apply to any edition of *The Pragmatic Programmer* ebook.

Thank you for your continued support. We hope to hear from you again soon!

The Pragmatic Bookshelf



SAVE 30%

with coupon code

BUYANOTHER2026

when you checkout at pragprog.com

From Objects to Functions


Build applications quicker and with less effort using functional programming and Kotlin. Learn by building a complete application, from gathering requirements to delivering a microservice architecture following functional programming principles. Learn how to implement CQRS and EventSourcing in a functional way to map the domain into code better and to keep the cost of change low for the whole application life cycle.

If you're curious about functional programming or you are struggling with how to put it into practice, this guide will help you increase your productivity composing small functions together instead of creating fat objects.

Uberto Barbini

(468 pages) ISBN: 9781680508451. \$47.95

<https://pragprog.com/book/uboop>




A Common-Sense Guide to AI Engineering

Want to build an LLM-powered app but don't know where to begin? Can't get past a proof-of-concept? With this step-by-step guide, you can master the underlying principles of AI engineering by building an LLM-powered app from the ground up. Tame unpredictable models with prompt and context engineering. Use evals to keep them on track. Give chatbots the knowledge to answer anything a user wants to know. Equip agents with the tools and smarts to actually get the job done. By the end, you'll have the intuition and the confidence to build on top of LLMs in the real world.

Jay Wengrow

(300 pages) ISBN: 9798888651933. \$59.95

<https://pragprog.com/book/jwpaieng>




Vector Search with JavaScript

Make search results smarter and more useful for everyday users and deliver more relevant results with vector search. Go beyond keyword matching to build search experiences that understand meaning, context, and similarity. Use AI-powered techniques to create recommendation systems, personalized search, and content discovery tools. Implement vector search from the ground up with step-by-step guidance, real-world examples, and hands-on coding. Generate embeddings, construct vector indexes, and optimize search accuracy with practical methods that integrate seamlessly into JavaScript applications. Whether refining an existing project or developing a new one, unlock the power of AI-driven search to create smarter, more intuitive user experiences.

Ben Greenberg

(128 pages) ISBN: 9798888651735. \$35.95

<https://pragprog.com/book/bgvector>




Machine Learning in Elixir

Stable Diffusion, ChatGPT, Whisper—these are just a few examples of incredible applications powered by developments in machine learning. Despite the ubiquity of machine learning applications running in production, there are only a few viable language choices for data science and machine learning tasks. Elixir's Nx project seeks to change that. With Nx, you can leverage the power of machine learning in your applications, using the battle-tested Erlang VM in a pragmatic language like Elixir. In this book, you'll learn how to leverage Elixir and the Nx ecosystem to solve real-world problems in computer vision, natural language processing, and more.

Sean Moriarity

(372 pages) ISBN: 9798888650349. \$61.95

<https://pragprog.com/book/smelixir>




Risk-First Software Development, Second Edition

Not all software projects go according to plan: many fail due to overlooked problems, misaligned stakeholders, or rigid methodologies. This book offers a groundbreaking framework for thinking differently by identifying risk at the center of every decision. You'll gain the vocabulary, tools, and confidence to identify, evaluate, and mitigate risks before they derail your project. Whether you're managing a startup product, steering an enterprise system, or trying to incorporate new technologies such as AI, risk-first helps you get your team aligned, spot trouble before it hits, and build software that delivers.

Rob Moffat

(265 pages) ISBN: 9798888651803. \$52.95

<https://pragprog.com/book/rmrfsd>




The Healthy Programmer, Second Edition

To keep doing what you love, you need to maintain your own systems, not just the ones you write code for. Regular exercise and proper nutrition help you learn, remember, concentrate, and be creative—skills critical to doing your job well. In this book you'll see how to change your work habits, master exercises that make working at a computer more comfortable, and develop a plan to keep fit, healthy, and sharp for years to come.

Joe Kutner

(242 pages) ISBN: 9798888651537. \$57.95

<https://pragprog.com/book/jkthp2>




A Common-Sense Guide to Data Structures and Algorithms in Python, Volume 2

Want to write code that pushes the boundaries of speed, space savings, and scalability? Then you need more advanced data structures and algorithms. Go beyond Big O notation and evaluate the true efficiency of each algorithm you design. Pull out data structures such as B-trees, bit vectors, and Bloom filters to wrangle big data. Wield techniques like caching, randomization, and fingerprinting to tame even the most demanding applications. With simple language, clear diagrams, and practice exercises and solutions, this book makes these topics easy to grasp. Go beyond the basics and use these next-level concepts to build software that's ready to take on the challenges of the real world.

Jay Wengrow

(500 pages) ISBN: 9798888651322. \$75.95

<https://pragprog.com/book/jwpython2>




Effective Haskell

Put the power of Haskell to work in your programs, learning from an engineer who uses Haskell daily to get practical work done efficiently. Leverage powerful features like Monad Transformers and Type Families to build useful applications. Realize the benefits of a pure functional language, like protecting your code from side effects. Manage concurrent processes fearlessly. Apply functional techniques to working with databases and building RESTful services. Don't get bogged down in theory, but learn to employ advanced programming concepts to solve real-world problems. Don't just learn the syntax, but dive deeply into Haskell as you build efficient, well-tested programs.

Rebecca Skinner

(668 pages) ISBN: 9781680509342. \$57.95

<https://pragprog.com/book/rshaskell>



The Pragmatic Bookshelf

The Pragmatic Bookshelf features books written by professional developers for professional developers. The titles continue the well-known Pragmatic Programmer style and continue to garner awards and rave reviews. As development gets more and more difficult, the Pragmatic Programmers will be there with more titles and products to help you stay on top of your game.

Visit Us Online

This Book's Home Page

<https://pragprog.com/book/ubaidev>

Source code from this book, errata, and other resources. Come give us feedback, too!

Keep Up-to-Date

<https://pragprog.com>

Join our announcement mailing list (low volume) or follow us on Twitter @pragprog for new titles, sales, coupons, hot tips, and more.

New and Noteworthy

<https://pragprog.com/news>

Check out the latest Pragmatic developments, new titles, and other offerings.

Buy the Book

If you liked this ebook, perhaps you'd like to have a paper copy of the book. Paperbacks are available from your local independent bookstore and wherever fine books are sold.

Contact Us

Online Orders: <https://pragprog.com/catalog>

Customer Service: support@pragprog.com

International Rights: translations@pragprog.com

Academic Use: academic@pragprog.com

Write for Us: <http://write-for-us.pragprog.com>